

Dependable Distributed Systems
Exercise week 2
October 6th, 2022

Exercise 1

With reference to the synchronization of physical clocks, provide the definition of internal and external clock synchronization. In addition, consider a system composed of two processes p_1 and p_2 and one UTC server ps . Let us assume that:

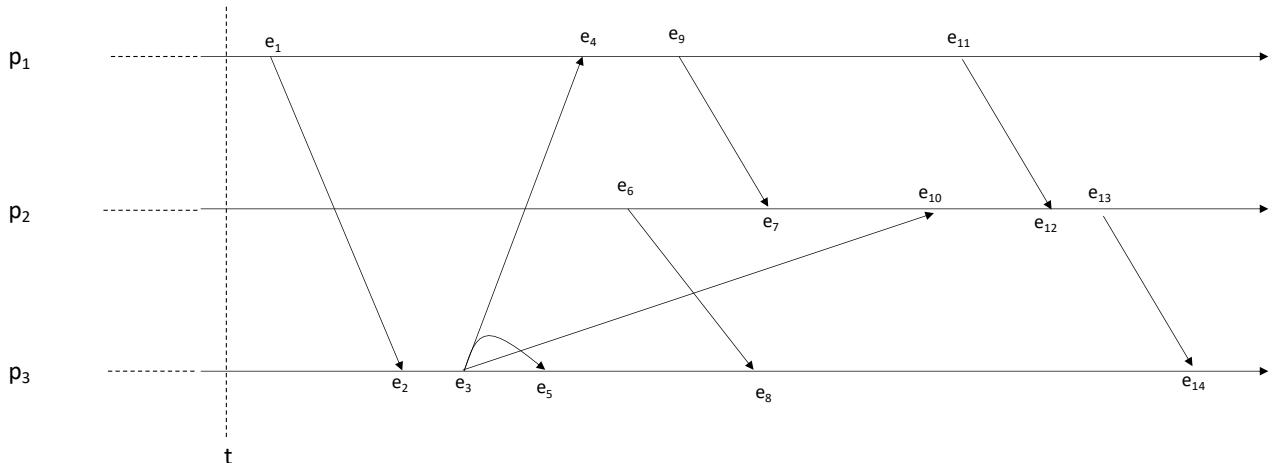
- p_1 and p_2 communicate with ps by using perfect point-to-point links
- the maximum latency of the channel between ps and p_1 is 1 ms
- the maximum latency of the channel between ps and p_2 is 2 ms.

In addition, let us assume that p_1 and p_2 start a clocks synchronization procedure at a certain time t by running the Christian algorithm. Answer to the following questions:

1. how much is the accuracy bound D_{ext} of the external synchronization obtained by p_1 and p_2 at the end of the synchronization?
2. Is the current system internally synchronized? If yes, determine the internal synchronization bound D_{int} obtained at the end of the procedure.

Exercise 2

Let us consider the execution history depicted in the figure

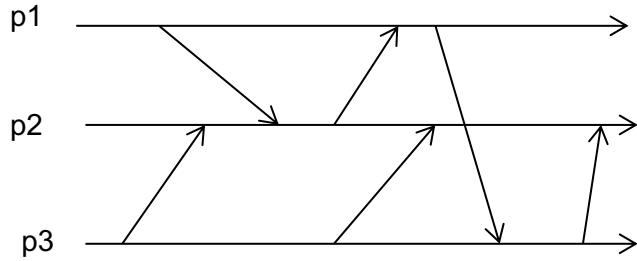


Given the run depicted in the figure state the truthfulness of the following sentences:			
a	According to the happened-before relation, $e_5 \rightarrow e_7$	T	F
b	According to the happened-before relation, $e_4 \parallel e_5$	T	F
c	Let CK_i be the variable storing the scalar logical clock of process p_i . Let us assume that at time t , $CK_i = 0$ for each process p_i . The logical clock CK_2 associated to e_6 is strictly larger than the logical clock CK_1 associated to e_1	T	F
d	Let CK_i be the variable storing the scalar logical clock of process p_i . If at time t $CK_1 = 0$ then the logical clock associated to e_8 is $CK_2=3$	T	F
e	Let CK_i be the variable storing the vector logical clock of process p_i . If at time t $CK_1 = [0, 0, 0]$ then the logical clock associated to e_8 is $CK_2=[3, 0, 3]$	T	F

For each point, provide a justification for your answer

Exercise 3

Describe timestamping techniques based on scalar logical clocks and vector logical clocks. In addition, considering the execution reported in Figure, answer to the following questions:



1. Apply the scalar clock timestamping technique to the execution assigning a timestamp to each event
2. Apply the vector clock timestamping technique to the execution assigning a timestamp to each event
3. List all pairs of concurrent events in the proposed execution

Exercise 4

Consider an asynchronous message passing system that uses vectors clock to implement some causal consistency check. The message passing system is composed by 4 processes with IDs 1 to 4, and, as usual, the ID is used as displacement in the vector clock (i.e., the locations are (p_1, p_2, p_3, p_4)).

Vector clocks are updated increasing before send. Processes communicate by point2point links. You start debugging process p_1 (the process with id 1) in the middle of the algorithm execution. You see the following stream of messages exiting and entering the ethernet card of process p_2 :

- Time 00:00 – EXITING: Send Message [MSG CONTENT] Vector Clock: $(1, 0, 0, 0)$
- Time 00:05 – ENTERING: Rcvd Message [MSG CONTENT] Vector Clock: $(1, 2, 3, 1)$

1. Draw an execution that justifies the vectors clocks you are seeing. Is such an execution unique?
2. There exists an execution that justifies the vector clocks and where there exists at least a process that does not send any message? Justify your answer.

Exercise 5

Let us consider a distributed system composed of N processes p_1, p_2, \dots, p_n each one having a unique integer identifier. Processes are arranged in line topology as in the following figure



Let us assume that there are no failures in the system (i.e., processes are always correct) and that topology links are implemented through perfect point-to-point links.

Write the pseudo-code of a distributed algorithm that is able to build the abstraction of a perfect point-to-point link between any pair of processes (also between those that are not directly connected).

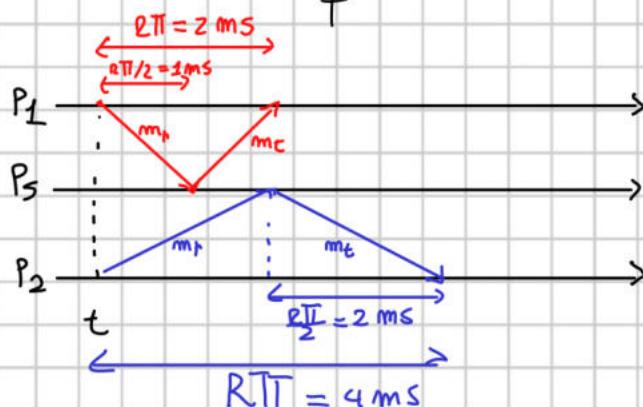
Exercise 1

With reference to the synchronization of physical clocks, we talk about **internal synchronization** when all the processes synchronize their clocks C_i between them and clocks are internally synchronized in a time interval Δ if $|C_i(t) - C_j(t)| < D$ for $i, j = 1 \dots N$ and for all time t in Δ , where $D > 0$ is the synchronization bound and $C_i - C_j$ the clocks at processes p_i and p_j . We have **external synchronization** when processes synchronize their clock C_i with an authoritative external source S (UTC) and clocks C_i are externally synchronized with a time source S if for each time interval Δ : $|S(t) - C_i(t)| < D$.

info:

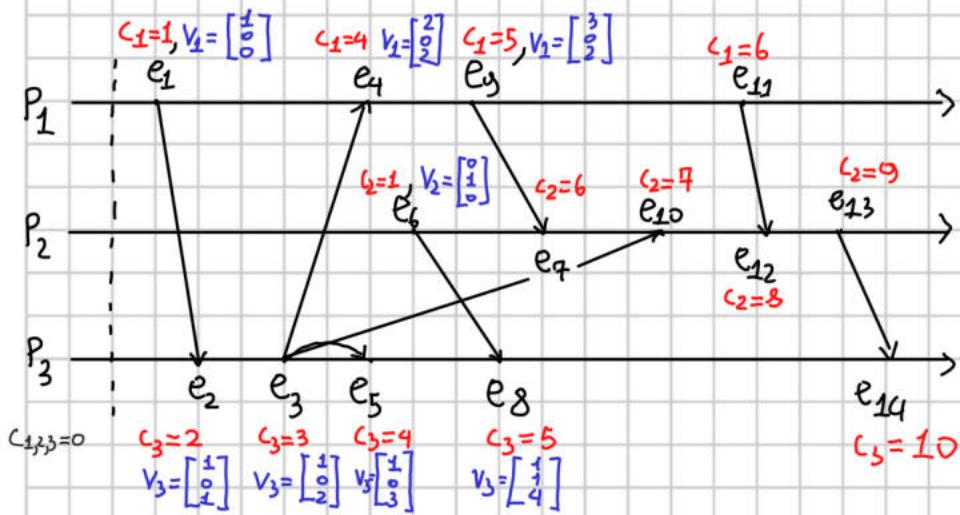
- there are P_1, P_2, P_3
- $P_1 \leftrightarrow P_3 \pm 1 \text{ ms}$
- $P_2 \leftrightarrow P_3 \pm 2 \text{ ms}$

Simulation of Christian algorithm:



- Accuracy bound D_{ext} of the external synchronization obtained by P_1 is $D_{ext_1} = \pm \frac{RTT_1}{2} = \pm 1 \text{ ms}$, while $D_{ext_2} = \pm \frac{RTT_2}{2} = \pm 2 \text{ ms}$
- Given that the system is externally synchronized is also internally synchronized in a bound $2D_{ext}$ that is $D_{int} = \pm 4 \text{ ms}$. We choose $\pm 4 \text{ ms}$ and not $\pm 2 \text{ ms}$ because we choose the greatest D_{ext} for consider the worst case scenario, that $\pm 2 \text{ ms}$ not cover. (Ex.: P_2 forward of D_2 ($+2 \text{ ms}$))

Exercise 2



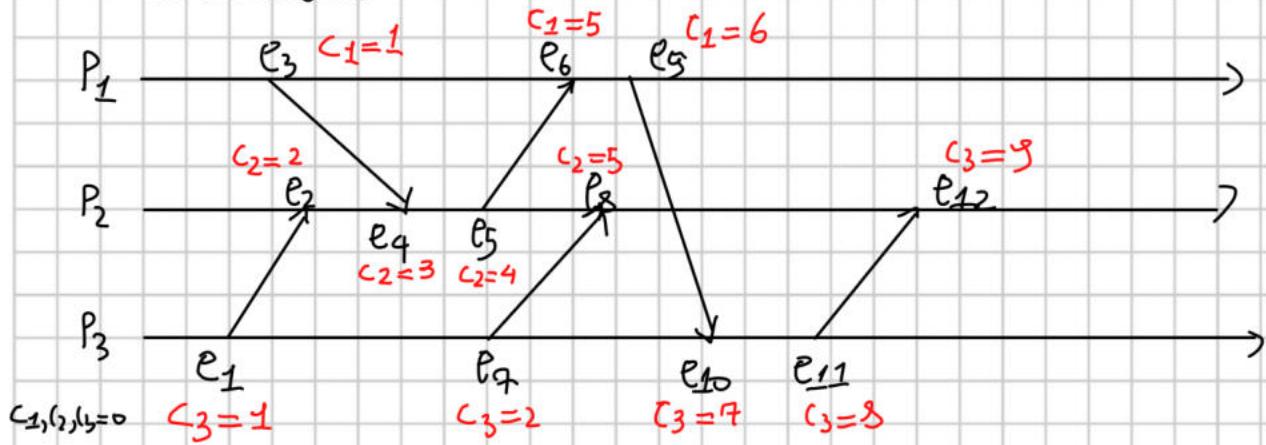
- False because the events $e_5 \parallel e_7$, not satisfy any of the 3 condition of the happened-before.
- True because v_1 associated to e_4 is $\begin{bmatrix} 2 \\ 0 \end{bmatrix}$ and v_3 associated to e_5 is $\begin{bmatrix} 1 \\ 3 \end{bmatrix}$ ($2 > 1$, $0 = 0$ but $2 < 3$)
- False because the timestamp of $e_1=1$ and timestamp of $e_6=1$, they are equal.
- False because $c_3=5$ associated to e_8 .
- False because $v_3=\begin{bmatrix} 1 \\ 4 \end{bmatrix}$ associated to e_8 .

Exercise 3

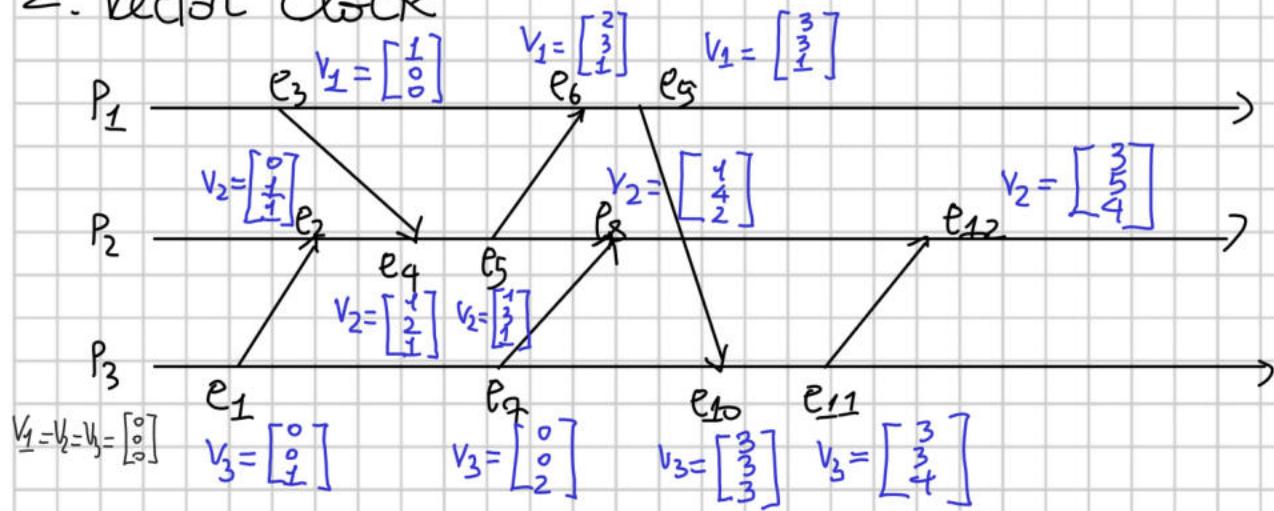
In **Scalar logical clocks** technique each process P_i initializes its logical clock $C_i=0$, P_i increase C_i of 1 when it generates an event. When P_i send a message m increase C_i of 1 and timestamp m with $t_s = C_i$, while when P_i receive a message m update $C_i = \max(t_s, C_i)$

In **Vector clocks** technique for a set of n processes the vector clock is composed by an array of n integers. Each process P_i maintain a vector clock V_i and timestamps events by mean of its vector clock.

1. Scalar clock



2. Vector clock



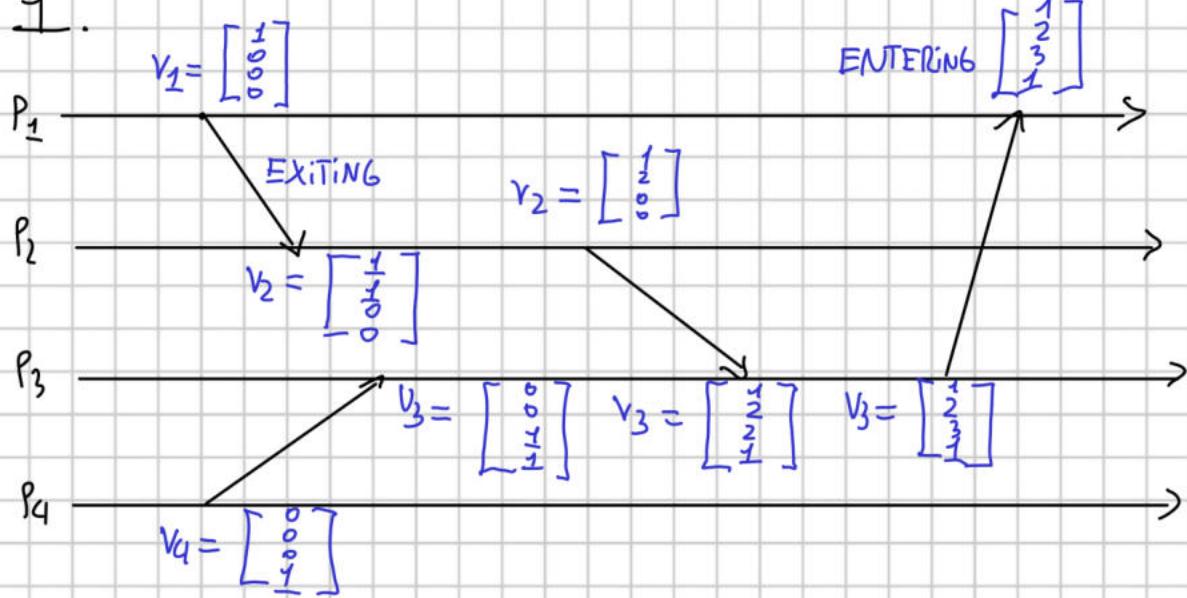
3. We can see the concurrent event using the vector clocks:

$e_1 \parallel e_3$, $e_2 \parallel e_3$, $e_2 \parallel e_7$, $e_5 \parallel e_7$, $e_4 \parallel e_7$, $e_3 \parallel e_4$,
 $e_9 \parallel e_7$, $e_6 \parallel e_7$, $e_8 \parallel e_5$, $e_8 \parallel e_{10}$, $e_8 \parallel e_{11}$

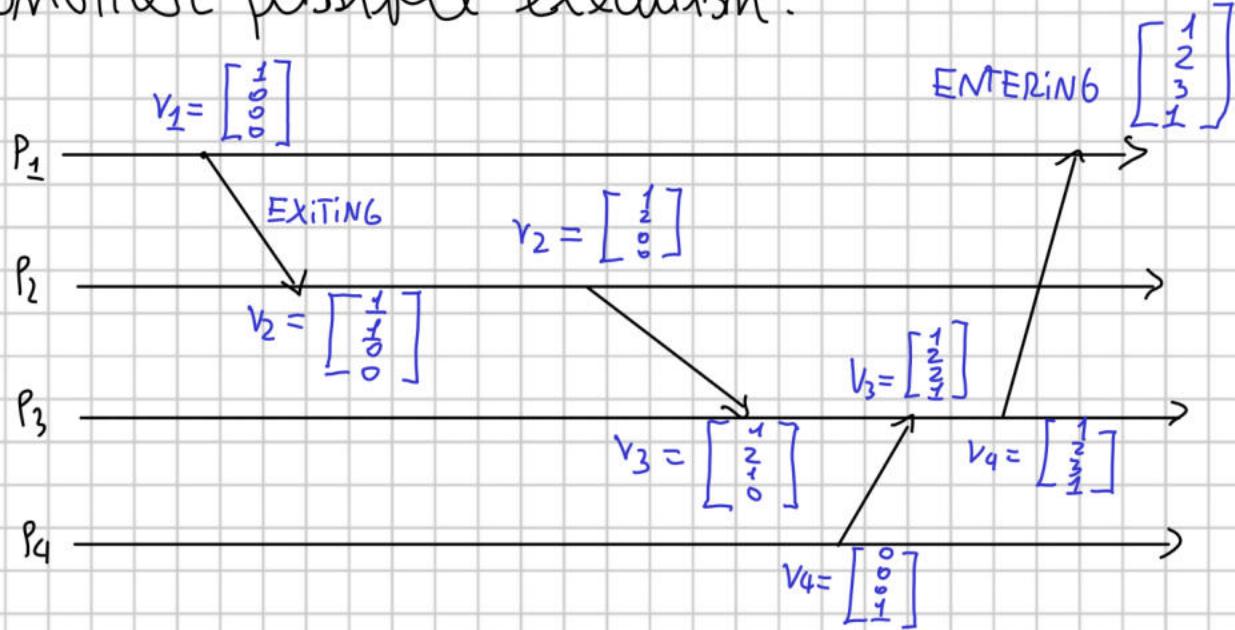
Exercise 4

4 processes: P_1, P_2, P_3, P_4 each with a vector clock V_i

1.

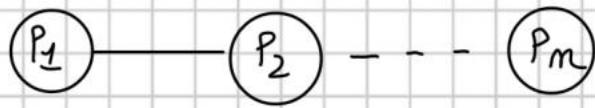


this execution is not unique because we show another possible execution:

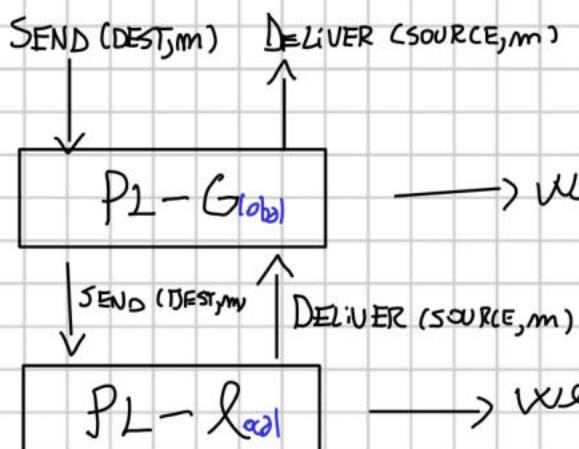


2. it's not possible that a process send any message, because in V_1 at time 00:05 enter a vector with all data of array at least 1 which implies that every process send at least a message.

Exercise 5



there are m processes, each with an identifier $i=1, \dots, N$.
each message is for another process $j=1, \dots, N$. Each
process can send the message only to the near process,
that are, for a process P_i , P_{i+1} and P_{i-1} . We are
using perfect link that guarantee that if the endpoints
are correct the message will received. we want implementing
a perfect line P2P given that each process sends
a message from the opposite side it delivered it.
we have two event handler, one for send and one
for deliver, we indicate with P the sending process
and with Q the receiver process.
these are the components and the events:



→ we want to implement this component

→ we have the perfect link component

algorithm:

implements: Perfect P2P PL-G

uses: P2P-link PL-L

every process have an ID unique.

Upon event < PL-G, SEND | q, m > do:
IF q > ID then ↳ component
 ↳ parameters

 trigger < PL-L, SEND | ID+1, < ID, q, m > >

else IF q < ID then

 trigger < PL-L, SEND | ID-1, < ID, q, m > >

else

 trigger < PL-L, SEND | ID, < ID, q, m > >

Upon event < PL-L, DELIVER | S, < p, q, m > do:

IF q == ID then

 trigger < PL-G, DELIVER | p, m >

else IF q > ID then

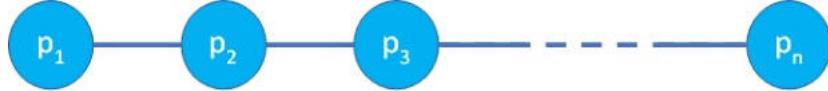
 trigger < PL-L, SEND | ID+1, < p, q, m > >

else

 trigger < PL-L, SEND | ID-1, < p, q, m > >

Exercise week 2 - October 6th, 2022 - Exercise 5. Commented Solution

Let us consider a distributed system composed of N processes p_1, p_2, \dots, p_n each one having a unique integer identifier. Processes are arranged in line topology as in the following figure



Let us assume that there are no failures in the system (i.e., processes are always correct) and that topology links are implemented through perfect point-to-point links.

Write the pseudo-code of a distributed algorithm that is able to build the abstraction of a perfect point-to-point link between any pair of processes (also between those that are not directly connected).

General Comments

The text says that processes have access to a perfect point-to-point link primitive, let us refer to this component with \mathcal{PL}_l . This primitive allows every process to exchange messages with at most two other processes, the neighbor processes in the line.

Let us recall the events and properties of a perfect point-to-point primitive.

Module:

Name: PerfectPointToPointLinks, **instance** *pl*.

Events:

Request: $\langle pl, Send \mid q, m \rangle$: Requests to send message *m* to process *q*.

Indication: $\langle pl, Deliver \mid p, m \rangle$: Delivers message *m* sent by process *p*.

Properties:

PL1: Reliable delivery: If a correct process *p* sends a message *m* to a correct process *q*, then *q* eventually delivers *m*.

PL2: No duplication: No message is delivered by a process more than once.

PL3: No creation: If some process *q* delivers a message *m* with sender *p*, then *m* was previously sent to *q* by process *p*.

Every process can thus trigger *Send* events and catch *Deliver* events of the \mathcal{PL}_l component.

The exercise requests to define a protocol a link abstraction, let us name it \mathcal{PL}_g , that implements a perfect point-to-point link between all pairs of processes. In the initial setting, every process can exchange messages only with at most two processes, namely the neighbor processes on the line. It follows that if two processes are not linked by \mathcal{PL}_l , such as processes p_1 and p_4 , they cannot exchange messages in the current setting. The target of the exercise is the definition of a protocol that implements a communication primitive guaranteeing all the properties of a perfect point-to-point link between all pairs of processes.

Protocol Idea

Each process is associated with an integer identifier and it is placed in a topology (the line) such that all the processes with a lower identifier are placed on one side (from the prospective of a selected process) and all processes with a higher identifier are located on the other side. Furthermore, the processes are placed in order on the line with respect to their identifiers. It follows that if process p_1 wants to communicate with process p_4 , it can send its messages to process p_2 , that will relay such messages to p_3 , which finally forwards them to p_4 . More in detail, if the destination of a message is a process with a higher identifier with respect to a selected process, then it is sufficient to forward the message to the neighbors with greater ID till reaching the destination process, and vice-versa for a destination with a lower identifier.

Pseudo-code

Algorithm 1 \mathcal{PL}_g

```

1: procedure INIT
2:   | ID  $\leftarrow$  unique integer identifier

3: upon event  $\langle \mathcal{PL}_g, Send \mid dest, m \rangle$  do
4:   | if dest  $>$  ID then
5:     |   trigger  $\langle \mathcal{PL}_l, Send \mid ID + 1, \langle ID, dest, m \rangle \rangle$ 
6:   | else if dest  $<$  ID then
7:     |   trigger  $\langle \mathcal{PL}_l, Send \mid ID - 1, \langle ID, dest, m \rangle \rangle$ 
8:   | else
9:     |   trigger  $\langle \mathcal{PL}_l, Send \mid ID, \langle ID, dest, m \rangle \rangle$ 

10: upon event  $\langle \mathcal{PL}_l, Deliver \mid q, \langle source, dest, m \rangle \rangle$  do
11:   | if dest == ID then
12:     |   trigger  $\langle \mathcal{PL}_g, Deliver \mid source, m \rangle$ 
13:   | else if dest  $>$  ID then
14:     |   trigger  $\langle \mathcal{PL}_l, Send \mid ID + 1, \langle source, dest, m \rangle \rangle$ 
15:   | else
16:     |   trigger  $\langle \mathcal{PL}_l, Send \mid ID - 1, \langle source, dest, m \rangle \rangle$ 

```

Pseudo-code Comments

The first upon block (lines 3-9) rules how a process must act when the *Send* operation of the global perfect point-to-point primitive we are developing, \mathcal{PL}_g , is triggered. It is a tiny procedure that simply starts the propagation of a new message, $\langle source, dest, m \rangle$, over the proper link, \mathcal{PL}_l . Again, a content m (let use this word to distinguish m from $\langle source, dest, m \rangle$) targeted to a process with an higher identifier ($dest$) is reachable (on the line) by relaying the content to the neighbor with higher ID and vice-versa.

The second upon block rules how a process that has received a message $\langle source, dest, m \rangle$ from the primitive \mathcal{PL}_l (namely the local perfect point-to-point link) must act. This procedure compares the *dest* field contained inside the received message with the process identifier. If the process is the destination of a content, then it triggers the *Deliver* operation of the primitive we are developing (content m was targeted to this process). Otherwise, it continues the propagation of the message $\langle source, dest, m \rangle$ over the proper perfect point-to-point link. Notice that we need to forward the information about the source ID *source* and the destination id *dest* insider the message, otherwise a process (not linked with the source process) receiving a content m from \mathcal{PL}_l cannot assert whether the content is targeted to it-self and which process was its source.

Informal Correctness Proofs

We briefly check the correctness of the provided pseudo-code. The \mathcal{PL}_g primitive we defined must handle two events, $\langle \mathcal{PL}_g, Send \mid dest, m \rangle$ and $\langle \mathcal{PL}_g, Deliver \mid source, m \rangle$. Furthermore, it must guarantee the *reliable delivery, no duplication* and *No creation* property of a perfect point-to-point link abstraction.

Reliable delivery: if a $\langle \mathcal{PL}_g, Send \mid dest, m \rangle$ event is generated (from the application layer), then the first upon procedure reacts to this event by starting the propagation of the message $\langle source, dest, m \rangle$ over the proper perfect point-to-point link, triggering a *Send* event on the \mathcal{PL}_l primitive. The \mathcal{PL}_l abstraction guarantees all the perfect point-to-point link properties between two linked processes. Therefore, every *Send* event of the \mathcal{PL}_l component generates a *Deliver* event on the destination process (e.g. a $\langle \mathcal{PL}_l, Send \mid p_2, \langle source, dest, m \rangle \rangle$ event on process p_1 generates the $\langle \mathcal{PL}_l, Deliver \mid p_1, \langle source, dest, m \rangle \rangle$ event on process p_2). The second upon procedure guarantees the propagation of the message $\langle source, dest, m \rangle$ till process $p_i, i = dest$. Only process $p_i, i = dest$ triggers the $\langle \mathcal{PL}_g, Deliver \mid source, m \rangle$ event. The propagation policy eventually guarantees the occurrence of this event.

No duplication: \mathcal{PL}_l guarantees no duplication of the messages diffused by the primitive. For a single $\langle \mathcal{PL}_g, Send \mid dest, m \rangle$ event a single message $\langle source, dest, m \rangle$ is generated and propagated over the line using \mathcal{PL}_l , furthermore every process relays at most once a single $\langle source, dest, m \rangle$ message.

No creation: $\langle \mathcal{PL}_g, Deliver \mid source, m \rangle$ event can be triggered only from the reception of a message $\langle source, dest, m \rangle$. The message $\langle source, dest, m \rangle$ is initially generated by the processes only managing a $\langle \mathcal{PL}_g, Send \mid dest, m \rangle$ event.

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

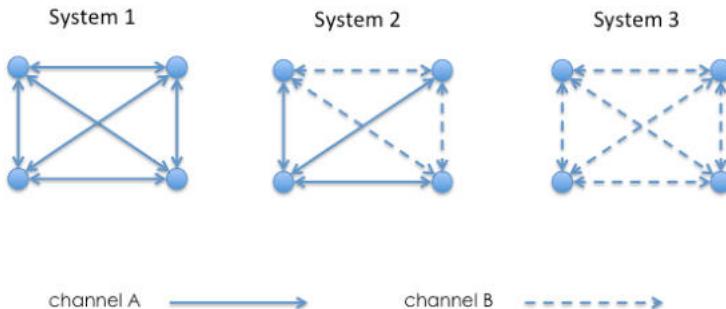
AA 2022/2023

Lecture 8 – Exercises

Ex 1: Let channel A and channel B be two different types of point-to-point channels satisfying the following properties:

- channel A: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j by time $t+\delta$.
- channel B: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j with probability p_{cons} ($p_{\text{cons}} < 1$).

Let us consider the following systems composed by 4 processes p_1, p_2, p_3 and p_4 connected through channels A and channels B.



Assuming that each process p_i is aware of the type of channel connecting it to any other process, answer to the following questions:

1. is it possible to design an algorithm implementing a perfect failure detector in system 2 if only processes having an outgoing channel of type B can fail by crash? true
2. is it possible to design an algorithm implementing a perfect failure detector in system 2 if any process can fail by crash? False
3. is it possible to design an algorithm implementing a perfect failure detector in system 3? False

For each point, if an algorithm exists write its pseudo-code, otherwise show the impossibility.

Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ that communicate by exchanging messages on top of a line topology, where p_1 and p_n are respectively the first and the last process of the network.

Initially, each process knows only its left neighbour and its right neighbour (if they exist) and stores the respective identifiers in two local variables LEFT and RIGHT.

Processes may fail by crashing, but they are equipped with a perfect oracle that notifies at each process the new neighbour (when one of the two fails) through the following primitives:

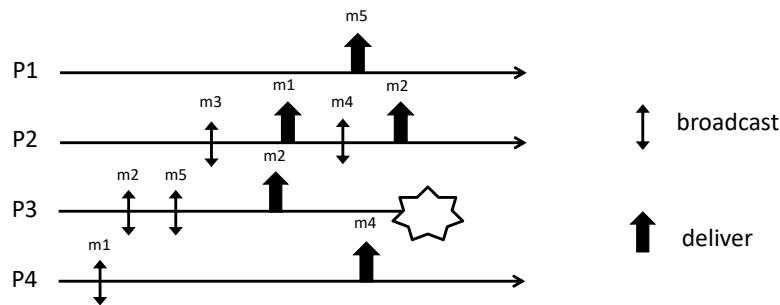
- **Left_neighbour(p_x):** process p_x is the new left neighbour of p_i
- **Right_neighbour(p_x):** process p_x is the new right neighbour of p_i

Both the events may return a NULL value in case p_i becomes the first or the last process of the line.

Each process can communicate only with its neighbours.

Write the pseudo-code of an algorithm implementing a Leader Election primitive assuming that channels connecting two neighbour processes are perfect.

Ex 3: Consider the partial execution depicted in the Figure



Answer to the following questions:

1. Complete the execution in order to have a run satisfying Uniform Reliable Broadcast.
2. Complete the execution in order to have a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast.
3. Complete the execution in order to have a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast.

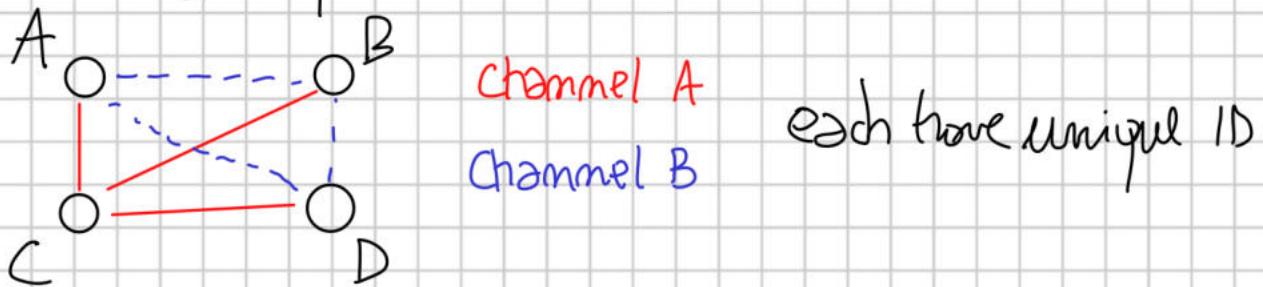
NOTE: In order to solve the exercise you can add broadcast, deliver and crash events but you cannot remove anything from the run.

Ex 4: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ identified through unique integer identifiers. Processes may communicate using perfect point-to-point links. Links are available for any pair of processes.

Processes may fail by crash and each process has access to a perfect failure detector. Modify the algorithms implementing a distributed mutual exclusion abstraction discussed during the lectures to allow them to tolerate crash failures.

Exercise 1

1. Yes is possible because all processes are linked with channel A that guarantee that the message sent by a process arrive to the destination.

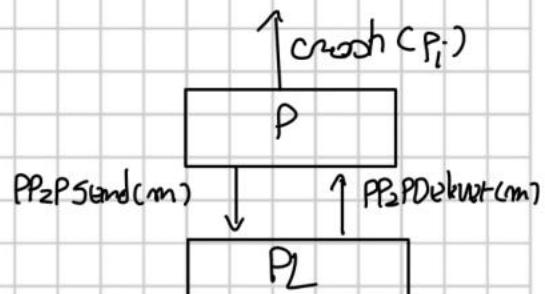


each have unique ID

Fundamental is C that can not crash, because otherwise we can't guarantee the specification of a perfect failure detector, in particular the accuracy. All messages of A, B and D must pass to C, we can write this pseudo code:

implements: PerfectFailureDetector, P

uses: PerfectP2PLinks, PL



upon event < P, init > do

alive := π ; $\{p_1, p_2, p_3, p_4\}$
detected := \emptyset ;

timer₁ = 5

starttimer(timer₁)

→ we use two timer

upon event timeout timer₁ do

forall $p_j \in \pi$ do

trigger < PL, send | LIST, alive_j > to p_j

starttimer(timer₂) → timer₂ is equal to timer₁

Upon event < PL, Deliver | LIST, I > from p_j

alive_j = alive_j \cup I

→ all processes merge alive list
of C because of channel A

When timeout timer₂ do

For every $p_j \in \Pi$ do

IF $p_j \notin \text{alive}$; AND $p_j \notin \text{detected}$;

$\text{detected}_j = \text{detected} \cup \{p_j\}$

trigger $\langle p_j, \text{crash} | p_j \rangle$

trigger $\langle p_L, \text{send} | \text{HEARTBEAT} \rangle$ to p_j

$\text{alive}_j = \emptyset$

$\text{start timer}(t_{\text{empty}})$

↗ search for crashed processes

Upon event $\langle p_L, \text{Deliver} | \text{HeartBeat} \rangle$ from p_j

$\text{alive}_j = \text{alive}_j \cup \{p_j\}$ → all update alive; list, in particular
c will receive all correct HB requests

in this way in the first timer₁ the processes A, B, C derive the
alive processes known by C, that is always correct because have chA.
in the second) timer with heartBeat C add who is alive to
the list and all work until C is the only that don't fail.

2. False, given that from exercise 1 we learn
that without the processes C this algorithm
don't work, if C crash the communication
between processes use channels A, that don't
guarantee accuracy of PerfectFailureDetector, the
system become asynchronous.

3. False the channels B are required for the
PerfectFailureDetector, system 3 have only channel A,
can't guarantee accuracy

Exercise 2

implements: LeaderElection, LE

Uses: Obstacle O, P2P link PL

Upon event $\langle LE, \text{init} \rangle$ do

suspected = \emptyset

leader = P_1 \rightarrow process at beginning at start is the leader

LEFT = get_left()

RIGHT = get_right

Upon event $\langle O, \text{LEFT_NEIGHBOUR } | P_x \rangle$ do

left = P_x

if left = null \rightarrow must be that leader failed, i am the new leader

leader = self

trigger $\langle PL, \text{PP2PSend } | \text{NEW_LEADER}, \text{leader} \rangle$ to right

Upon event $\langle O, \text{RIGHT_NEIGHBOUR } | P_x \rangle$ do

right = P_x \rightarrow we simply update, the problem is on left

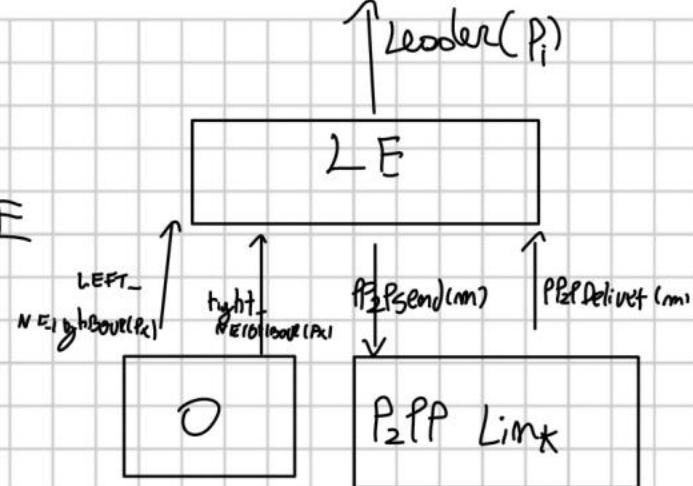
Upon event $\langle PL, \text{PP2P Deliver} | \text{NEW_LEADER}, () \rangle$ from left

leader = ()

trigger $\langle LE, \text{leader} | \text{leader} \rangle$

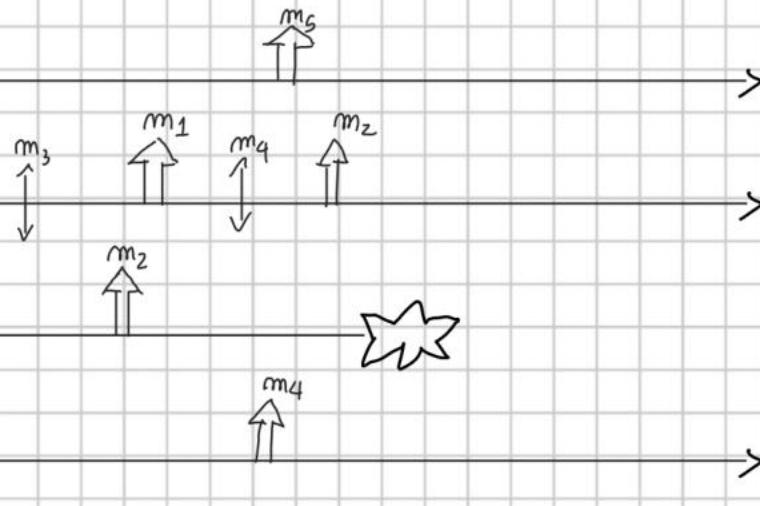
trigger $\langle PL, \text{PP2PSend} | \text{NEW_LEADER}, \text{leader} \rangle$ to right

idea is to set initial leader the process at start, when receive the left neighbour from obstacle we check that is null, that implies leader is crashed, and in the case add that process as new leader and propagate the news.



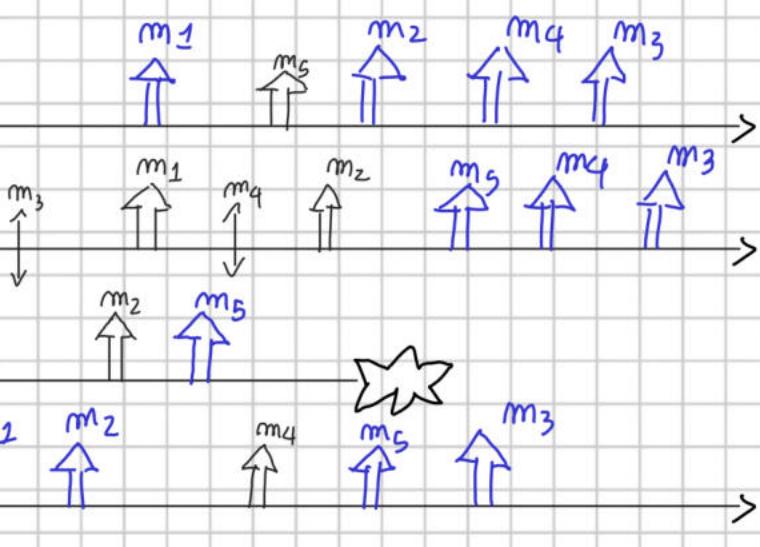
Exercise 3

P₁



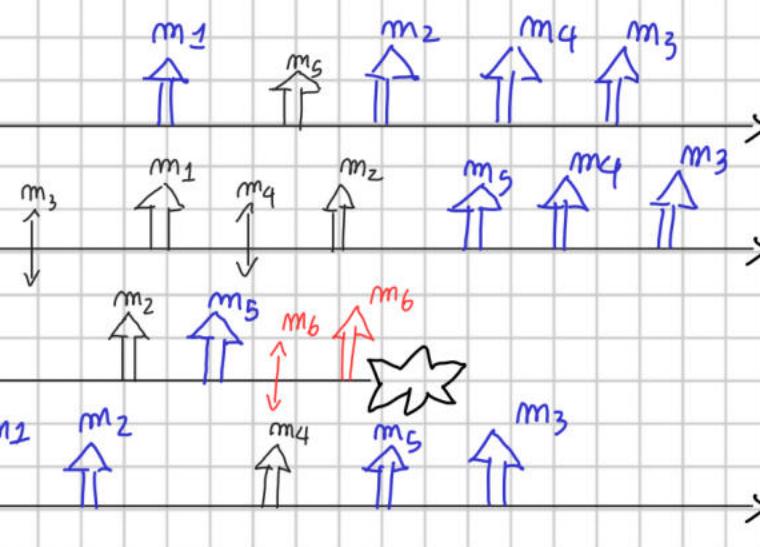
1.

P₁



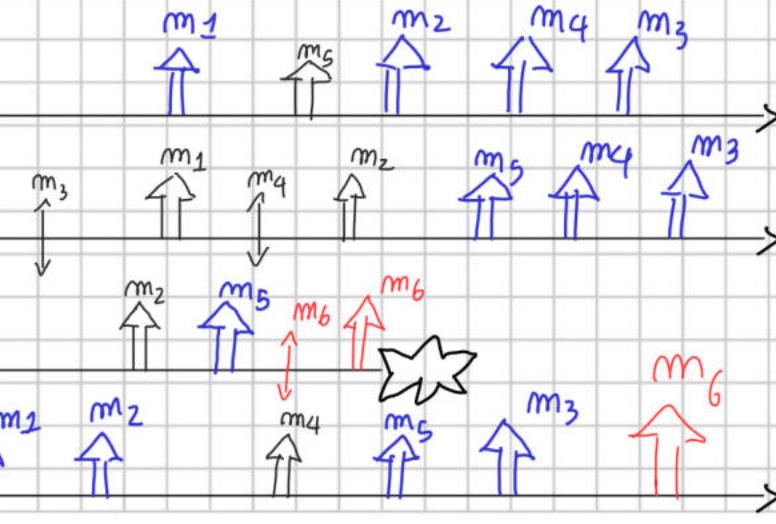
2.

P₁



3.

P₁



Exercise 4

We modify the Rigid-Agarwal's Algorithm for mutual exclusion such that tolerate crash failures. Given that processes uses P2P Links, all processes are linked each other, and have a Perfect Failure Detector.

uses: P2P Links, P
PerfectFailureDetector, P

init:

replies = 0
state = NCS
 $Q = \emptyset$
Last_Req = 0
Num = 0

begin

1. state = Requesting;
2. Num = num + 1; Last_Req = num;
3. $\forall i = 1, \dots, N$ send REQUEST(num) to p_i , $i \neq ID$
4. wait until replies = m - 1
5. state = CS
6. CS
7. $\forall t \in Q$ send REPLY to t
8. $Q = \emptyset$; state = NCS; replies = 0

upon receipt REQUEST(t) from p_j

IF state = CS or (state = Requesting and $\{Last_Req, i\} < \{t, j\}$)
insert in $Q \{t, j\}$

else

Send REPLY to p_j
num = max(t, num)

upon receipt REPLY from p_j

replies = replies + 1

Upon receipt Crash(p) \rightsquigarrow PerfectFailureDetector notify that p crashed

IF p in Q
remove p from Q

m = m - 1 \rightsquigarrow we have minus processes

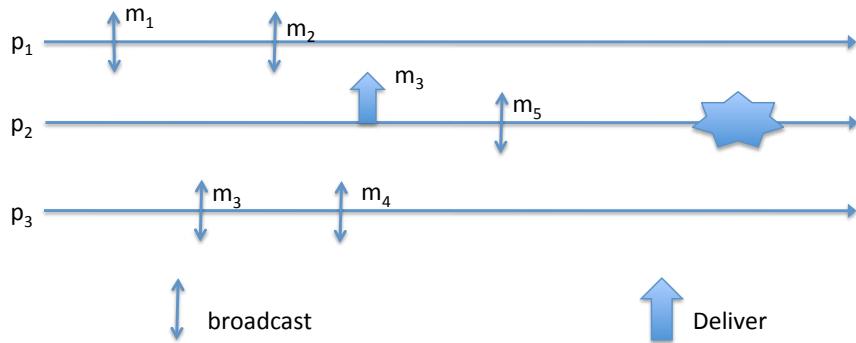
Simply we manage the crash event in the way that if crashed process is in Queue for enter in CS then we delete it. Also we decrease m as the process don't wait for another reply for enter in CS.

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 11 – Exercises
October 20th, 2022

Ex 1: Consider the partial execution depicted in the following figure:



1. Complete the execution in order to obtain a run satisfying *Best Effort Broadcast* but *not Reliable Broadcast*.
2. Complete the execution in order to obtain a run satisfying *Regular Reliable Broadcast* but *not Uniform Reliable Broadcast*.
3. Complete the execution in order to obtain a run satisfying *Uniform Reliable Broadcast*.

Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$. Each process is connected to all the others through fair-loss point-to-point links and has access to a perfect failure detector.

Write the pseudo-code of an algorithm implementing a Uniform Reliable Broadcast primitive.

Additionally, answer to the following questions:

1. Is it possible to provide a quiescent implementation of the Uniform Reliable Broadcast primitive?
2. Given the system model described here, is it possible to provide an implementation that uses only data structure with finite size?

Ex 3: Consider a distributed system composed by N servers $\{s_1, s_2, \dots, s_n\}$ and M clients $\{c_1, c_2, \dots, c_m\}$.

Each client c_i runs its algorithm and it can request to servers the execution of a particular task T_i . Servers will execute the task T_i and, after that, a notification will be sent to c_i that T_i has been completed.

The Figure shows the code executed by a generic client c_i .

Operation executeTask (T_i)	Upon pp2pdeliver (TASK_COMPLETED, T_i) from s_j
1. For each $s_i \in \{s_1, s_2, \dots, s_n\}$ 2. pp2psend (TASK_REQ, T_i, c_i) to s_i ;	1. trigger completedTask (T_i);

Write the pseudo-code of an algorithm, executed by servers, able to allocate tasks assuming that:

- Once clients ask for a task execution, they remain blocked until the task is not terminated.
- Any two clients c_i and c_j can concurrently require the execution of two different tasks T_i and T_j ;
- Each task is univocally identified by the pair (T_i, c_i) ;
- Each server can manage at most one task at every time;
- At most $N-1$ servers can crash;
- Servers can use a uniform consensus primitive;
- Servers can use a failure detector P ;
- Servers communicate through a uniform reliable broadcast primitive.

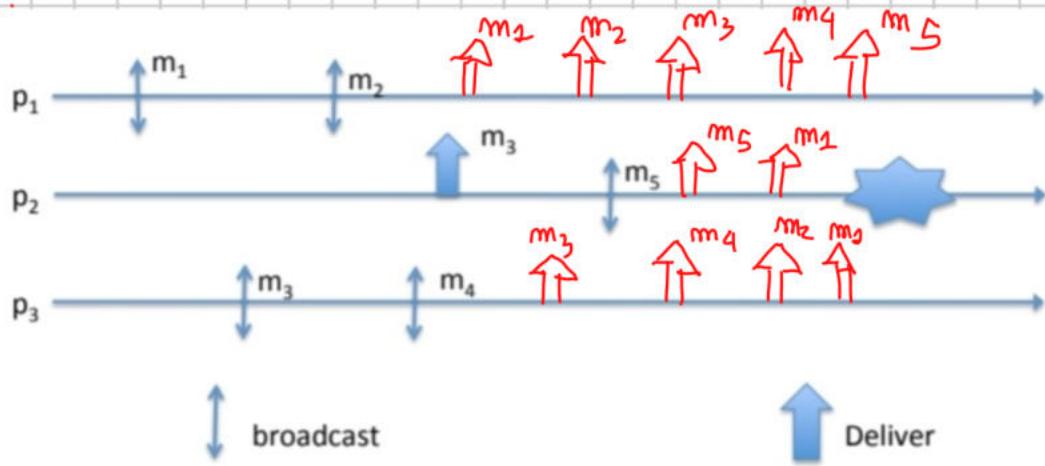
Note that, if a server crashes while executing a task, such task needs to be re-allocated and re-processed by a different server.

Ex 4: Consider a distributed system formed by n processes p_1, p_2, \dots, p_n connected along a ring i.e., a process p_i is initially connected to a process $p_{(i+1) \bmod n}$ through a unidirectional perfect point-to-point link.

Write the pseudo-code of a distributed algorithm implementing a consensus primitive.

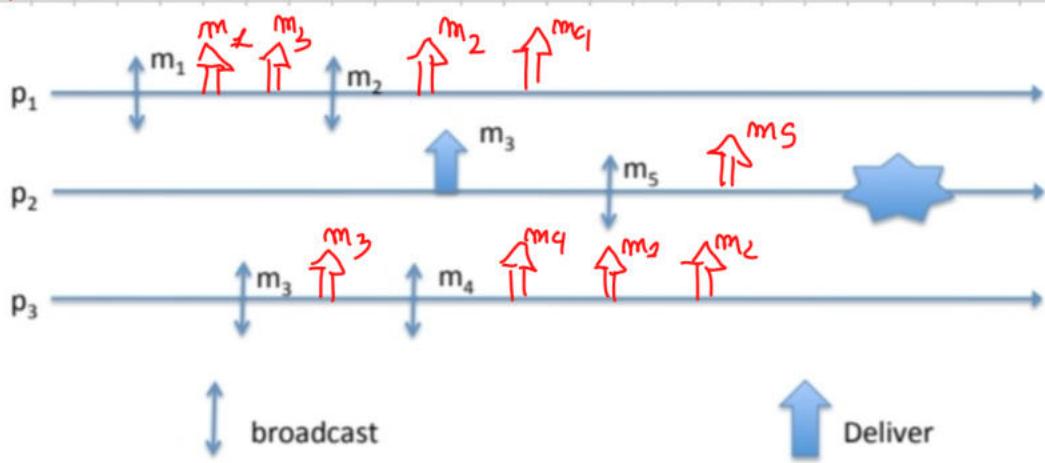
Exercise 1

1.



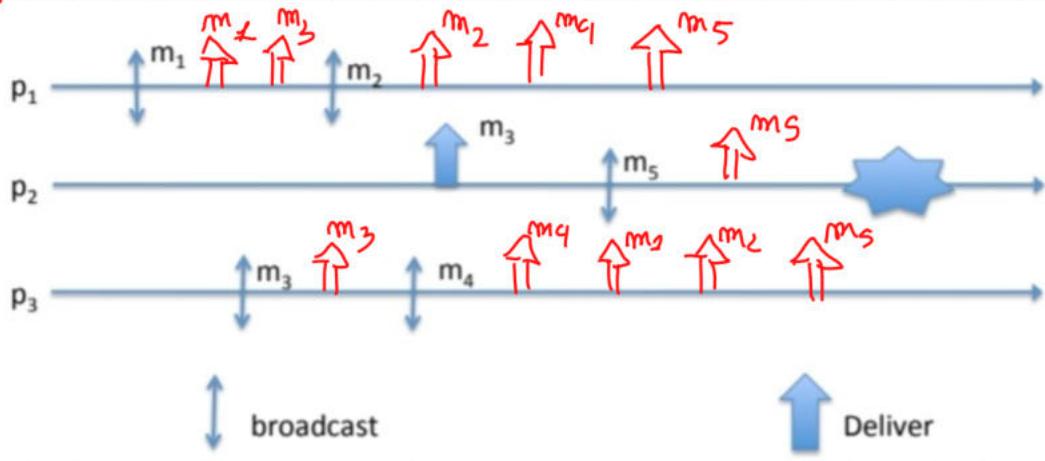
Best effort
Broadcast

2.



Reliable
Broadcast

3.



Uniform
Reliable
Broadcast

Exercise 2

implements Uniform Reliable Broadcast, URB

Uses: fair-loss-point-to-point, FL

Perfect Failure Detector, P

init:

$$t_{met} = \Delta, \text{ start}(t_{met})$$

$$\text{delivered} = \emptyset$$

$$\text{pending} = \emptyset$$

$$\text{correct} = \Pi$$

$$\forall m \text{ do } \text{ACK}[m] = \emptyset$$

Upon event URB-Broadcast(m) do

$$\text{pending} = \text{pending} \cup \{(self, m)\}$$

$\forall p_i \in \text{correct}$ do trigger $\text{FLsend}(m)$ to p_i

Upon event $\text{FLDeliver}(m)$ from p_j

$$\text{ACK}[m] = \text{ACK}[m] \cup \{p_j\}$$

$\text{if } (p_j, m) \notin \text{pending} \text{ and } m \notin \text{delivered}$

$$\text{pending} = \text{pending} \cup \{(p_j, m)\}$$

$\forall p_i \in \text{correct}$ do trigger $\text{FLsend}(m)$ to p_i

Upon event $\text{Crash}(p)$ do

$$\text{correct} = \text{correct} \setminus \{p\}$$

Function $\text{canDeliver}(m)$ returns Boolean is

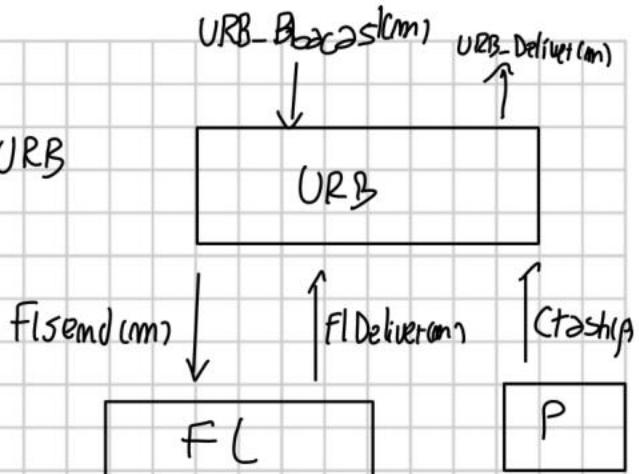
$\text{return } (\text{correct} \subseteq \text{ACK}[m]),$

Upon exist $(s, m) \in \text{pending}$ such that $\text{canDeliver}(m)$

$\wedge m \notin \text{delivered}$ do

$$\text{delivered} = \text{delivered} \cup \{m\}$$

trigger $\text{URB-Deliver}(m)$

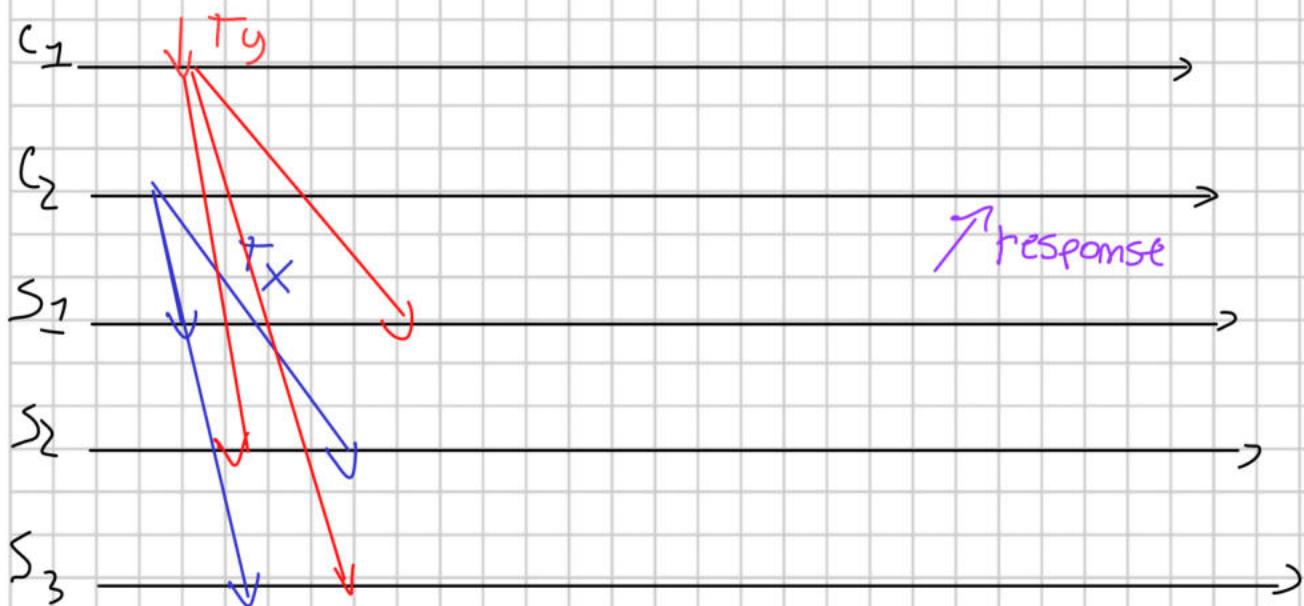
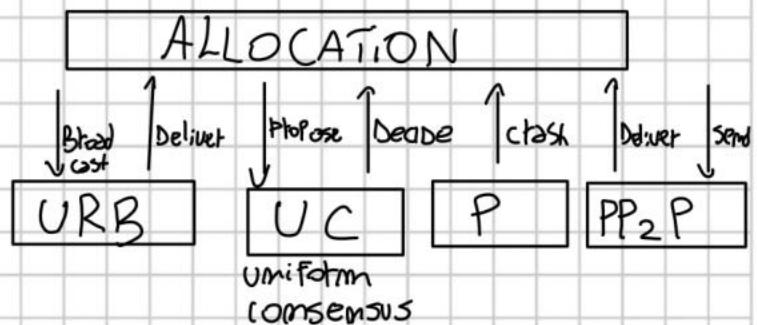
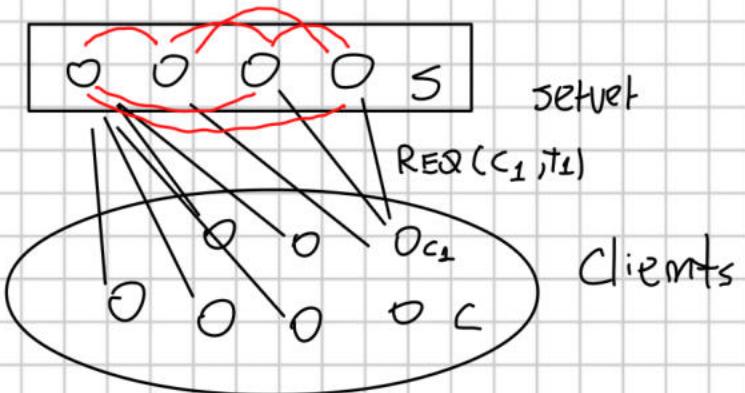


Upon timer expired do
for all $p \in \text{pending}$ and $p \notin \text{delivered}$ do
 trigger $f\{\text{send}(m)\}$ to p
start timer (timer)

1 it is possible to have a quiescent implementation of P2P link sending the ACK, avoid to transmitting forever

2. if the system continuously broadcast the messages, you can not avoid to use an infinite data structure, if at some point you wait them you can bound the memory. You can erase something only when you are sure no more messages concerning the delivered messages.

Exercise 3



init:

$\text{pending} = \emptyset$

$\text{busy} = \text{false}$

$\text{current_allocation} = \text{null}$

Upon event $\langle PL, PP_2P \text{ Deliver} | TASK_REQ, T_i, C_i \rangle$ from C_i
 if $\langle T_i, C_i \rangle$ is not in $\text{current_allocation}$ do
 $\text{Pending} = \text{pending} \cup \{ \langle T_i, C_i \rangle \}$

When pending is not empty do

IF not busy do

$\text{candidate} = \text{select_from}(\text{pending})$

trigger $\langle UC, \text{Propose} | \langle self, \text{candidate} \rangle \rangle$

else

trigger $\langle UC, \text{Propose} | \langle self, \text{null} \rangle \rangle$

Upon event <UC, Decide | <id, task>>

current_allocation = current_allocation U {<id, task>}

IF id == self and task != null do

busy = true

execute(T_i)

busy = false

trigger <PL, PP, PSEND | TASK_COMPLETED, T_i> to c_i

trigger <URB, URBtoadcast | TASK_COMPLETED, T_i, c_i>

Upon event <URB, URBDeliver | TASK_COMPLETED, T_i, c_i>

current_allocation = current_allocation \ {<T_i, c_i>}

Pending = pending \ {<T_i, c_i>}

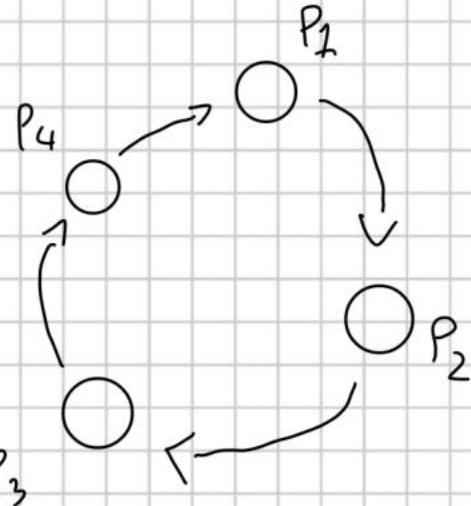
Upon event <P, ctash | s>

IF there exist <id, task> in current_allocation such that id == s do

current_allocation = current_allocation \ {<id, task>}

Pending = Pending U {Task}

Exercise 4



Use a token as a collector, you add your proposal and pass to the next, when come back to you all proposal are collected!
TOKEN initially is a empty list

initial

state = wait

next = $P_{(i+1) \bmod N}$

IF self = p_i

value proposed by p_i

TOKEN = TOKEN $\cup \{v\}$

trigger PP2PSend(TOKEN) to next

Upon event PP2PDeliver(TOKEN)

if size(TOKEN) == N \rightarrow all proposal received

decision = min(TOKEN) \rightsquigarrow min value of TOKEN is -1

trigger PP2PSend(decision) to next

trigger Decide(decision)

else

TOKEN = TOKEN $\cup \{v\}$

trigger PP2PSend(TOKEN) to next

Upon event PP2PDeliver(decision)

if state == wait

state = done

trigger PP2PSend(decision) to next

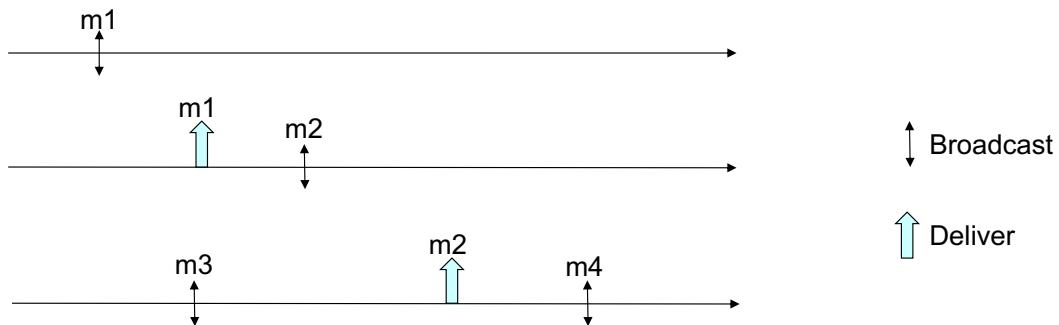
trigger Decide(decision)

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

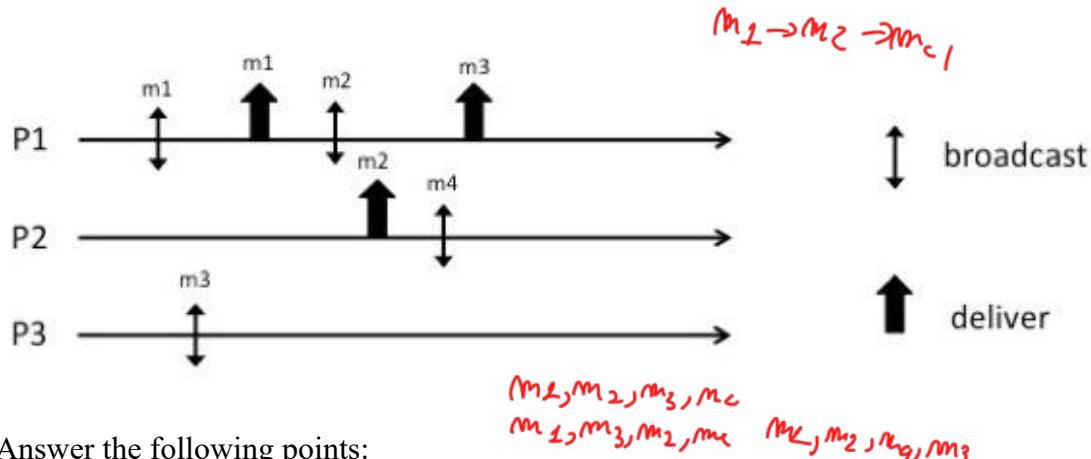
AA 2022/2023

Lecture 13 – Exercises
October 26th, 2022

Ex 1: Given the partial execution in Figure, provide all the delivery sequences such that both total order and causal order are satisfied



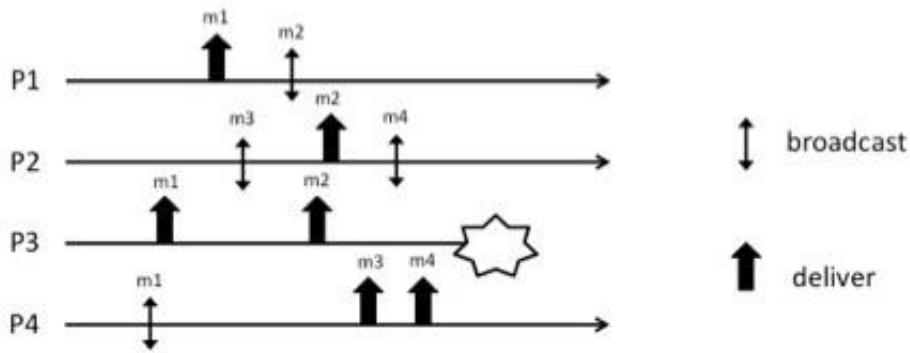
Ex 2: Let us consider the following partial execution



Answer the following points:

1. Provide all the possible sequences satisfying Causal Order
2. Complete the execution to have a run satisfying FIFO order but not causal order

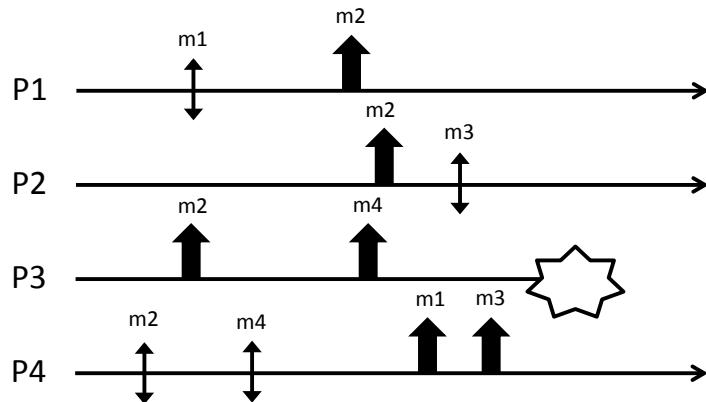
Ex 3: Let us consider the following partial execution



Answer the following points:

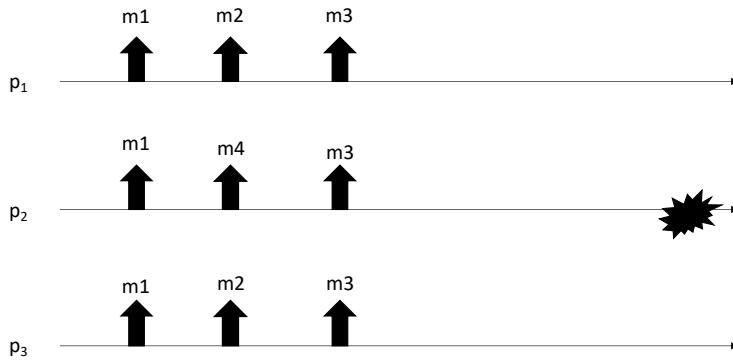
1. Provide the list of all the possible delivery sequences that satisfy both Total Order and Causal Order
2. Complete the history (by adding the missing delivery events) to satisfy Total Order but not Causal Order
3. Complete the history (by adding the missing delivery events) to satisfy FIFO Order but not Causal Order nor Total Order

Ex 4: Consider the message pattern shown in the Figure below and answer to the following questions:



1. Complete the execution in order to have a run satisfying Reliable Broadcast but not Uniform Reliable Broadcast.
2. Provide all the delivery sequences satisfying causal order and total order.
3. Provide all the delivery sequences violating causal order and satisfying TO(UA, WNUTO) but not satisfying TO(UA, SUTO)

Ex 5: Consider the partial execution in the following figure

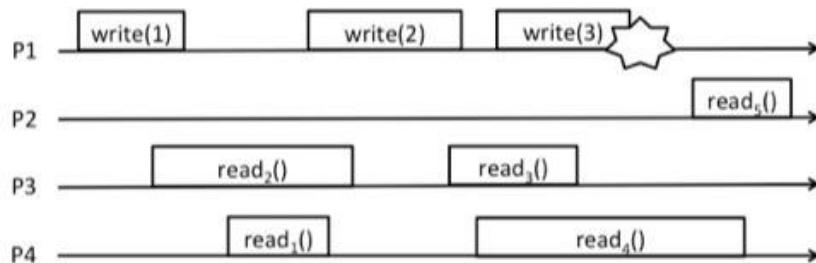


Given the run depicted in the figure state the truthfulness of the following sentences:

		T	F
a	The strongest agreement property satisfied is UA	T	F
b	The NUA agreement property is violated	T	F
c	The strongest ordering property satisfied is SUTO	T	F
d	The WUTO ordering property is satisfied	T	F
e	The SNUTO ordering property is violated	T	F
f	Let us assume we can add only one more delivery to p ₁ and p ₃ , it is not possible to get a run satisfying TO(NUA, SUTO)	T	F
g	If p ₂ is not going to deliver m ₄ then the strongest specification satisfied by the resulting execution is TO(UA, SUTO)	T	F
h	Let us assume we can add only one more delivery to p ₁ and p ₃ , it is possible to get a run satisfying TO(UA, WNUTO) but not satisfying TO(UA, WUTO)	T	F
i	If p ₂ is not faulty, the NUA agreement property is satisfied	T	F
j	If p ₂ is not faulty, the SUTO ordering property is satisfied	T	F

For each point, provide a justification for your answer

Ex 6: Consider the execution depicted in the following figure and answer the questions



1. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.
2. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.

Ex 7: Consider a distributed systems composed by a set of n processes p_1, p_2, \dots, p_n . Processes have a unique identifier and are structured as a binary tree topology. Messages are exchanged between processes over the edges of the tree which act like perfect point-to-point links. Each process p_i has stored the identifiers of its neighbors

into the local variables FATHER, R_CHILD e L_CHILD representing respectively the father of p_i , the right child and the left child (if they exists).

Assuming that processes are not going to fail, write the pseudo-code of an algorithm satisfying the following specification:

Events:

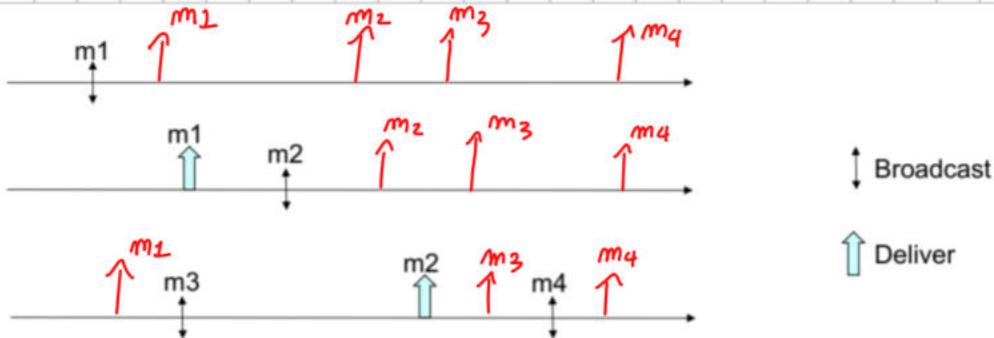
- **Request:** $\langle tob, \text{Broadcast} \mid m \rangle$: Broadcasts a message m to all processes.
- **Indication:** $\langle tob, \text{Deliver} \mid p, m \rangle$: Delivers a message m broadcast by process p .

Properties:

- *Validity*: If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement*: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- *Total order*: Let m_1 and m_2 be any two messages and suppose p and q are any two correct processes that deliver m_1 and m_2 . If p delivers m_1 before m_2 , then q delivers m_1 before m_2 .

Exercise 1

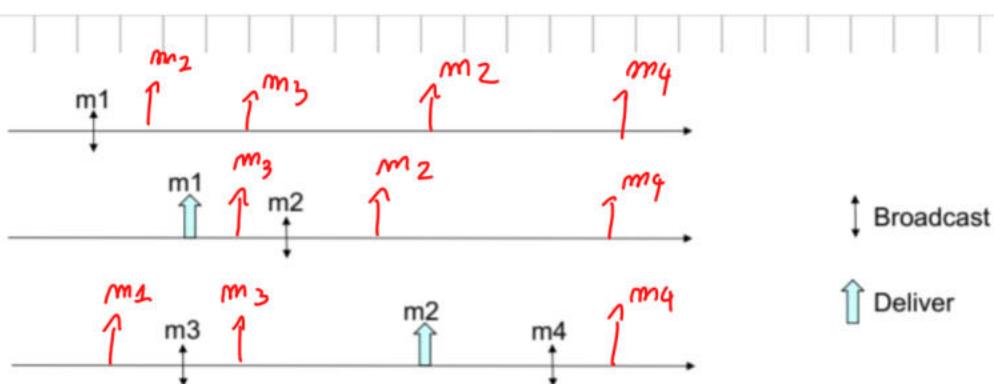
$m_1 \rightarrow m_2$ and $m_2 \rightarrow m_4$ (local order) $m_3 \rightarrow m_4$ (FIFO order)



m_1, m_2, m_3, m_4

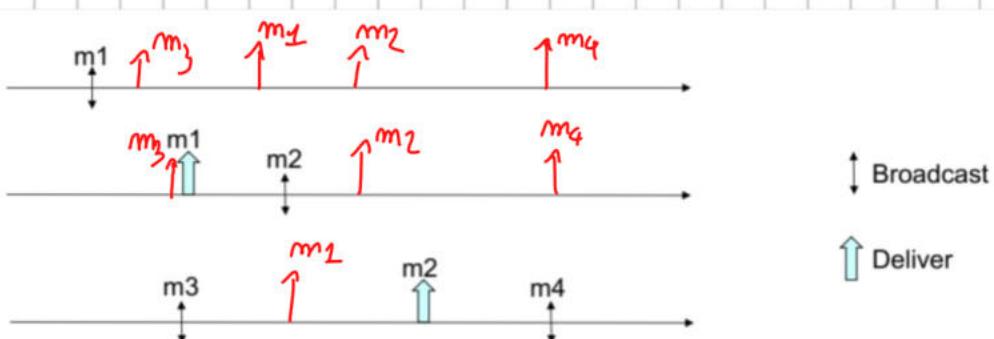
m_1, m_3, m_2, m_4

m_3, m_1, m_2, m_4



m_1, m_3, m_2, m_4

m_3, m_1, m_2, m_4



m_3, m_1, m_2, m_4

Exercise 2

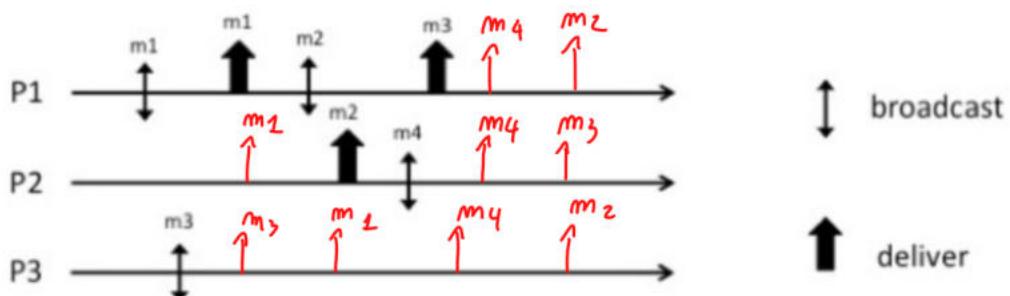
$m_1 \rightarrow m_2 \rightarrow m_4$ (local order + FIFO order)

1. sequence of deliveries:

We have four case:

- (m_1, m_2, m_3, m_4)
- (m_1, m_3, m_2, m_4)
- (m_1, m_2, m_4, m_3)
- (m_3, m_1, m_2, m_4)

2.



FIFO Broadcast but not causal

Exercise 3

Causal order = FIFO order + local order : $m_1 \rightarrow m_2 \rightarrow m_3$
 $m_3 \rightarrow m_4$

1. Total order and causal order

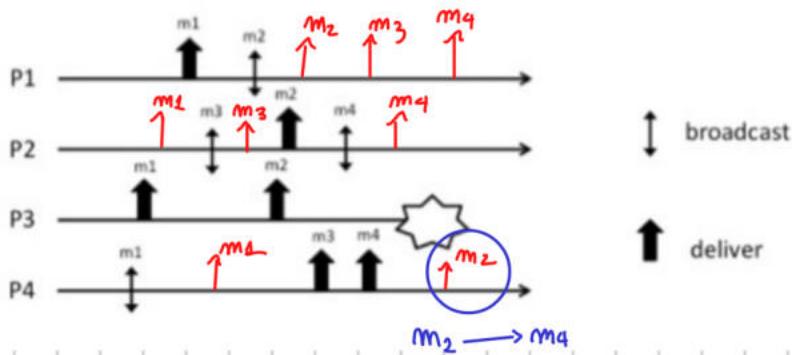
- m_1, m_2, m_3, m_4
- m_1, m_3, m_2, m_4

because of delivery of m_1 in P_1 is before of broadcast of m_3 in P_2 we can't have a sequence that start with m_3

2.

not exist an execution that guarantee Total order but not causal order given the delivery and odd only the missing delivery

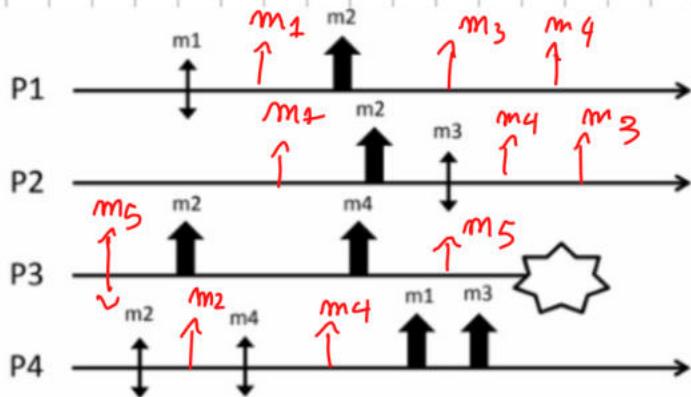
3.



FIFO order
but not
causal or
total

Exercise 4

1. Reliable broadcast but not URB



2. Causal order and total order

$$m_2 \rightarrow m_3 \text{ and } m_2 \rightarrow m_4$$

- m_1, m_2, m_3, m_4
- m_1, m_2, m_4, m_3
- m_2, m_1, m_5, m_4
- m_2, m_1, m_4, m_3
- m_2, m_4, m_1, m_3

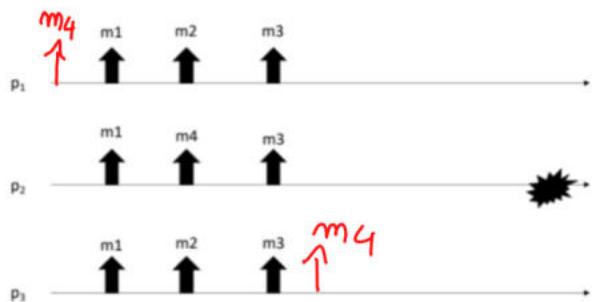
3. Violating causal order, TO (UA, WNUTO) but not satisfying TO (UA, SUTO)

- m_1, m_3, m_2, m_4
- m_1, m_3, m_4, m_2
- m_4, m_1, m_2, m_3
- m_4, m_2, m_1, m_3

↳ only the correct processes in P_3 we don't add any delay and any combination violate SUTO

Exercise 5

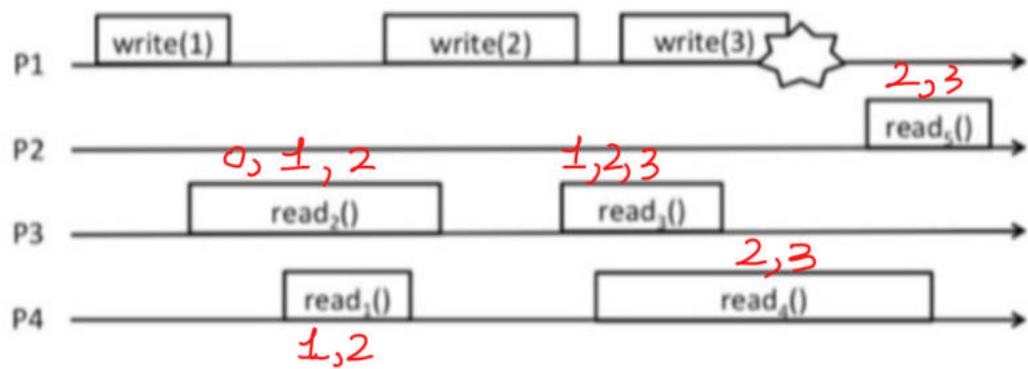
- a. F, we have NVA all correct processes have the same set of delivery, but p_2 faulty have m_4
- b. F, all the correct processes have the same set of deliveries, NVA is satisfied
- c. F, because p_2 deliver m_3 but not m_2 can't violate SUTO
- d. V, because p_1 and p_2 have the same order and p_2 respect the order of m_2 and m_3 , is UWTO
- e. F, this run satisfy UWTO that implies SNUTO
- f. V, because is missing m_2 in p_2 that not satisfy SUTO
- g. F, because like f p_2 deliver m_3 must deliver m_2 else for satisfy SUTO
- h. T, this run:



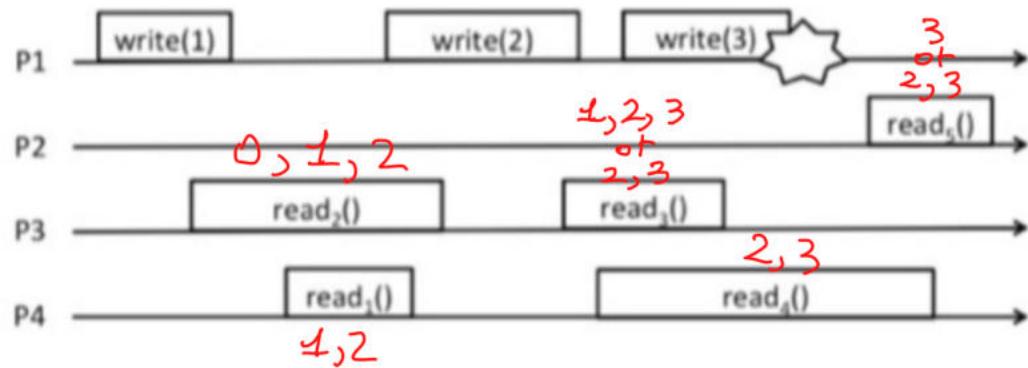
- i. F, because p_2 is correct and deliver m_4 but other correct don't deliver it, NVA is violated
- j. F, p_2 is missing m_2 and there is m_4 before m_3 .

Exercise 6

1. Rx possible value for Regular Register



2. Rx possible value for Atomic Register

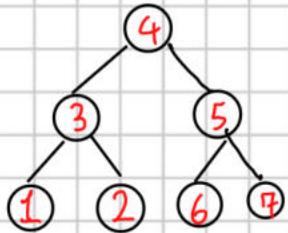


↳ because R₂ and R₃ are concurrent can have causal value

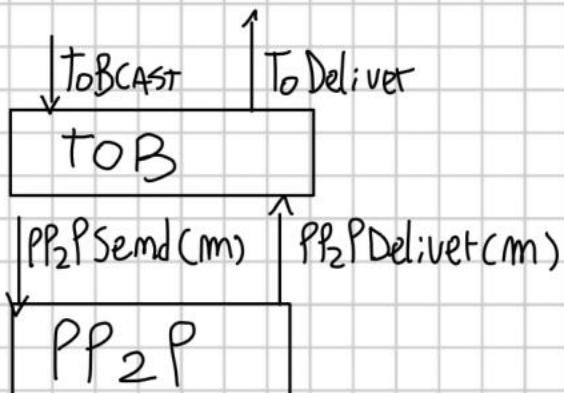
• R₃ depend on value chosen by R₁, if R₁ choose 2 then R₃ must return 2, 3 otherwise can return 1, 2, 3

• R₅ depend on value chosen by R₃, if R₃ choose 3 then R₅ must return 3 otherwise can return 2, 3

Exercise 7 - Solution



- processors have unique names.
- each process can talk only with near processes.
- links are P2P link.



$$P_1: \text{FATHER}_1 = P_3$$

$$\text{L_CHILD}_1 = \perp$$

$$\text{R_CHILD}_1 = \perp$$

One single process have FATHER; = \perp , here P_4 .

init:

FATHER; = get_father()

L_CHILD; = get_left()

R_CHILD; = get_right()

Sequence_number; = SM; = 0

LAST_Delivered; = 0

pending; = \emptyset

Upon event ToBcast(m)

if FATHER; $\neq \perp$

→ type of message

trigger PP2PSend('UNORDERED', m) to FATHER;

else

SM; = SM; + 1

trigger PP2PSend('ORDERED', m, SM;) to LEFT;

trigger PP2PSend('ORDERED', m, SM;) to RIGHT;

Upon event PP₂PDeliver ('UNORDERED', m) from P_j

If FATHER_i ≠ L

trigger PP₂PSend ('UNORDERED', m) to FATHER_i

else

SM_i = SM_i + 1

trigger PP₂PSend ('ORDERED', m, SM_i) to LEFT_i and RIGHT_i

Upon event PP₂PDeliver ('ORDERED', m, s)

pending_i = pending_i ∪ {<m, s>}

trigger PP₂PSend ('ORDERED', m, s) to LEFT_i and RIGHT_i

Upon $\exists \langle m, s_m \rangle \in \text{pending}_i$ s.t. $s_m = \text{LAST-DELIV.} + 1$

pending_i = pending_i / {<m, s>}

trigger ToDeliver(m)

LAST-DELIVERED = SM_i

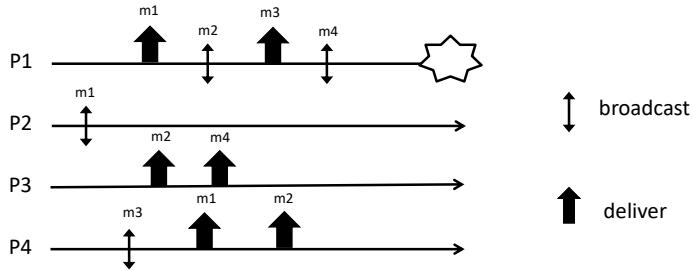
NO ID PROCESS

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 18 – Exercises
November 9th, 2022

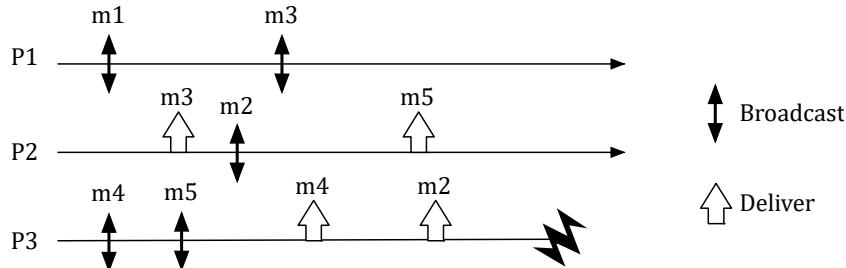
Ex 1: Consider the execution depicted in the Figure



Answer to the following questions:

1. Provide all the delivery sequences that satisfy both causal order and total order
2. Complete the execution in order to have a run satisfying TO(UA, WNUTO), FIFO order but not causal order

Ex 2: Consider the partial execution shown in the Figure and answer to the following questions:



1. Complete the execution in order to have a run satisfying the Regular Reliable Broadcast specification but not Uniform Reliable Broadcast one.
2. For each process, provide ALL the delivery sequences satisfying FIFO Reliable Broadcast but not satisfying causal order.
3. For each process, provide ALL the delivery sequences satisfying total order and causal order.

Ex 3: Consider a distributed system composed of N processes p_1, p_2, \dots, p_N , each having a unique identifier $myID$. Initially, all processes are correct (i.e. $correct = \{p_1, p_2, \dots, p_N\}$). Consider the following algorithm:

```

upon event Xbroadcast (m)
    mysn = mysn+1;
     $\forall p \in correct$ 
        pp2pSend ("MSG", m, mysn, myId);

upon event pp2pReceive ("MSG", m, sn, i)
    mysn= mysn+1;
    if (m  $\notin$  delivered)
        trigger XDeliver (m);
        delivered = delivered  $\cup$  {m};

upon event crash (pi)
    correct = correct / {pi}

```

Let us assume that: (i) links are perfect, (ii) the failure detector is perfect and (iii) initially local variables are initialized as follows $mysn=0$ e $delivered = \emptyset$.

Answer to the following questions:

1. Does the `Xbroadcast()` primitive implement a Reliable Broadcast, a Best Effort Broadcast or none of the two?
2. Considering only the ordering property of broadcast communication primitives discussed during the lectures (FIFO, Causal, Total), explain which ones can be satisfied by the `Xbroadcast()` implementation.

Provide examples to justify your answers.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through FIFO perfect point-to-point links and are structured through a line (i.e., each process p_i can exchange messages only with processes p_{i-1} and p_{i+1} when they exists). Processes may crash and each process is equipped with a perfect oracle (having the interface $new_right(p)$ and $new_left(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Perfect failure detector primitive.

Ex 5: Consider a distributed system composed of n processes p_1, p_2, \dots, p_n connected through a ring topology. Initially, each process knows the list of correct processes and maintains locally a `next` variable where it stores the id of the following process in the ring.

Each process can communicate only with its next trough FIFO perfect point-to-point channels (i.e. the process whose id is stored in the `next` variable).

Processes may fail by crash and each process has access to a perfect failure detector.

Write the pseudo-code of a distributed algorithm implementing a $(1, N)$ atomic register.

Ex 6: A transient failure is a failure that affects a process temporarily and that alter randomly the state of the process (i.e., when the process is affected by a transient failure, its local variables assume a random value).

Let us consider a distributed system composed by N processes where f_c processes can fail by crash and f_t processes can suffer transient failures between time t_0 and t_{stab} .

Let us consider the following algorithm implementing the Regular Reliable Broadcast specification

Algorithm 3.2: Lazy Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb ;

PerfectFailureDetector, **instance** \mathcal{P} .

```

upon event ⟨  $rb$ , Init ⟩ do
     $correct := \Pi$ ;
     $from[p] := [\emptyset]^N$ ;

upon event ⟨  $rb$ , Broadcast |  $m$  ⟩ do
    trigger ⟨  $beb$ , Broadcast | [DATA, self,  $m$ ] ⟩;

upon event ⟨  $beb$ , Deliver |  $p$ , [DATA,  $s$ ,  $m$ ] ⟩ do
    if  $m \notin from[s]$  then
        trigger ⟨  $rb$ , Deliver |  $s$ ,  $m$  ⟩;
         $from[s] := from[s] \cup \{m\}$ ;
        if  $s \notin correct$  then
            trigger ⟨  $beb$ , Broadcast | [DATA,  $s$ ,  $m$ ] ⟩;

upon event ⟨  $\mathcal{P}$ , Crash |  $p$  ⟩ do
     $correct := correct \setminus \{p\}$ ;
    forall  $m \in from[p]$  do
        trigger ⟨  $beb$ , Broadcast | [DATA,  $p$ ,  $m$ ] ⟩;

```

Answer to the following questions:

1. For every property of the Regular Reliable Broadcast specification, discuss if it guaranteed between time t_0 and t_{stab} and provide a motivation for your answer.
2. For every property of the Regular Reliable Broadcast specification, discuss if it eventually guaranteed after t_{stab} and provide a motivation for your answer.
3. Assuming that the system is synchronous, explain how you can modify the algorithm (no pseudo-code required) to guarantee that No Duplication, Validity and Agreement properties will be eventually guaranteed after t_{stab} .

Ex 7: A service is delivered through 3 components that sequentially handles the requests and it receives around 75 requests per second. The components, A, B and C, serve respectively 100, 80 and 90 requests per second. The clients are complaining about the response time of the system. There is the possibility either 1) to double one component employing a perfect load balancer or 2) to substitute one component with an improved one serving 40% more requests per second.

Which component upgrade and option reduce the response time the most?
Compute the expected response time.

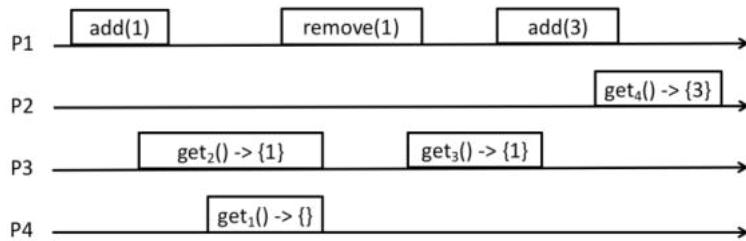
Ex 8: Consider a set object that can be accessed by a set of processes. Processes may invoke the following operations on the object:

- `add(v)`: it adds the value v in to the set
- `remove(v)` it removes the value v from the set
- `get()`: it returns the content of the set.

Informally, every `get()` operation returns all the values that have been added before its invocation and that have not been removed by any `remove()`.

For the sake of simplicity, assume that a value can be added/removed just once in the execution.

Consider the distributed execution depicted in the Figure



Answer to the following questions:

1. Is the proposed execution linearizable? Motivate your answer with examples.
2. Consider now the following property: “every `get()` operation returns all the values that have been added before its invocation and that have not been removed by any `remove()`. If an `add(v)/remove(v)` operation is concurrent with the `get`, the value v may or may be not returned by the `get()`”. Provide an execution that satisfy `get` validity and that is not linerizable.

Exercise 1

1. From the figure we have this CAUSAL ORDER:

$$m_1 \rightarrow m_2 \rightarrow m_4, m_3 \rightarrow m_4$$

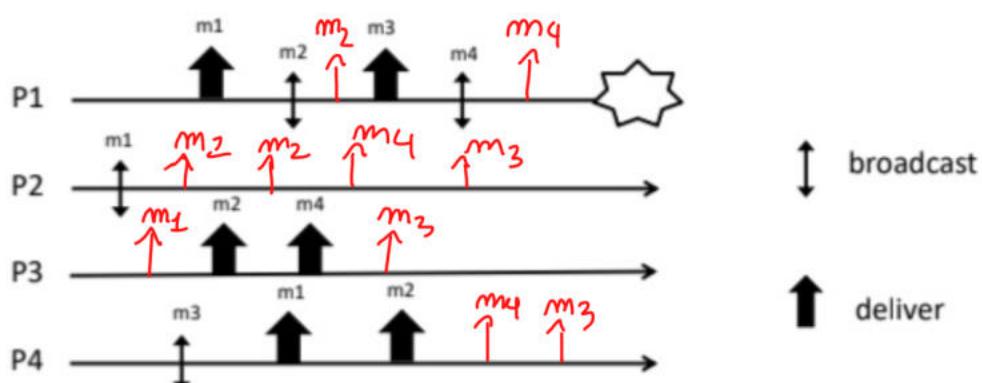
All possible deliveries that satisfy causal order and Total order:
for total order we must notice that m_3 must be delivered after m_1 (P_1), also m_2 must be delivered before m_4 (P_2)

$$m_1, m_2, m_3, m_4$$

$$m_1, m_3, m_2, m_4$$

(P_1 is faulty connect differently, but satisfies causal order)

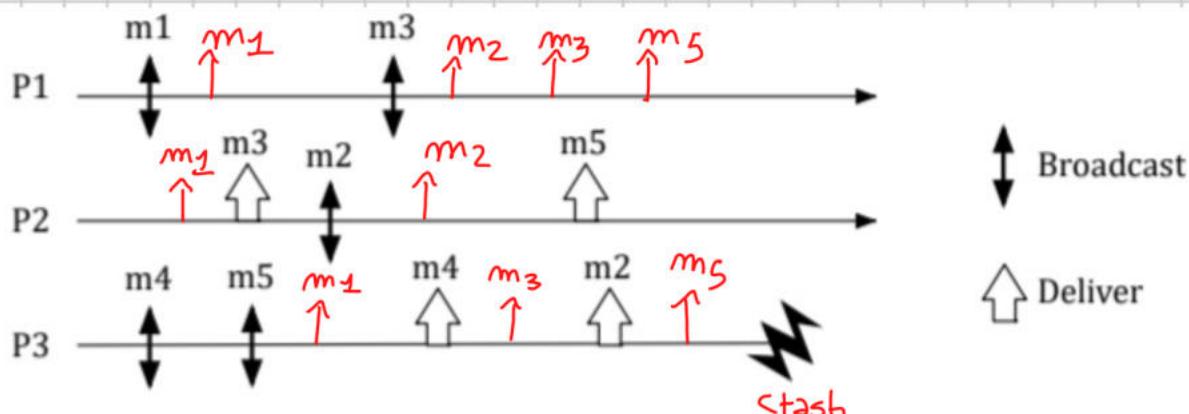
2. TO (UA, UNUTO), FIFO ORDER but not CAUSAL ORDER



FIFO ORDER: $m_2 \rightarrow m_4$ $m_4 \rightarrow m_3$

Exercise 2

1. RRB but not URB



P_3 faulty, m_4 is not delivered by P_3 it's on P_2 , not URB

2. FIFO : $m_4 \rightarrow m_5$, $m_1 \rightarrow m_3$

CAUSAL : $m_3 \rightarrow m_2$

All delivery satisfying FIFO Reliable Broadcast but not causal order for each process:

We must deliver also m_4 because we deliver m_5 for have FIFO order, we satisfy URB that implies RRB!

P_1 : $(m_1, m_2, m_3, m_4, m_5)$, $(m_1, m_2, m_4, m_3, m_5)$
 $(m_1, m_2, m_4, m_5, m_3)$, $(m_2, m_1, m_3, m_4, m_5)$
 $(m_2, m_1, m_4, m_3, m_5)$, $(m_2, m_1, m_4, m_5, m_3)$
 $(m_4, m_1, m_2, m_3, m_5)$, $(m_4, m_1, m_2, m_3, m_5)$
 $(m_4, m_1, m_2, m_5, m_3)$, $(m_4, m_2, m_1, m_3, m_5)$
 $(m_4, m_2, m_1, m_3, m_5)$, $(m_4, m_2, m_1, m_5, m_3)$
 $(m_1, m_4, m_2, m_3, m_5)$, $(m_1, m_4, m_2, m_5, m_3)$
 $(m_1, m_4, m_5, m_2, m_3)$, $(m_2, m_4, m_1, m_5, m_3)$
 $(m_2, m_4, m_5, m_1, m_3)$

P_2 : if we deliver m_4 before m_2 we violate the causal order in P_3 , for the other delivery we can have only combination but satisfy $m_4 \rightarrow m_5$ and $m_1 \rightarrow m_3$
EX.: $(m_4, m_2, m_3, m_1, m_5)$

P_3 : if in P_2 deliver m_4 before m_2 , we can have any combination, only respecting FIFO ORDER. EX. $(m_1, m_2, m_4, m_5, m_3)$

3. $m_4 \rightarrow m_5$ and $m_1 \rightarrow m_3 \rightarrow m_2$, in P_3 m_4 before m_2
 and in P_2 m_3 before m_5
- (m_4, m_1, m_3, m_2, m_5)
 - (m_1, m_4, m_3, m_2, m_5)
 - (m_1, m_3, m_4, m_2, m_5)
 - (m_4, m_1, m_3, m_5, m_2)
 - (m_1, m_4, m_5, m_3, m_2)

For each process there are possible delivery.

Exercise 3

N processes : $\{P_1, \dots, P_N\}$

upon event $X_{broadcast}(m)$

$$mySm = mySm + 1$$

$\forall p \in \text{correct}$

$P_{P_2} P_{send} ("MSG", m, mySm, myId)$

upon event $P_{P_2} P_{Receive} ("MSG", m, sm, i)$

$$mySm = mySm + 1$$

if ($m \notin \text{delivered}$)

$t_{trigger} X_{Deliver}(m);$

$$\text{delivered} = \text{delivered} \cup \{m\}$$

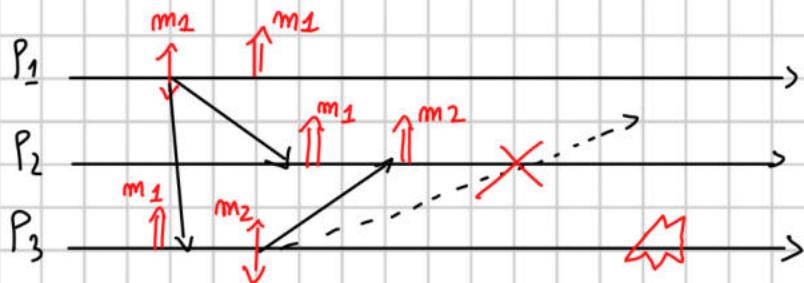
upon event $(\text{crash}(p_i))$

$$\text{correct} = \text{correct} / \{p_i\}$$

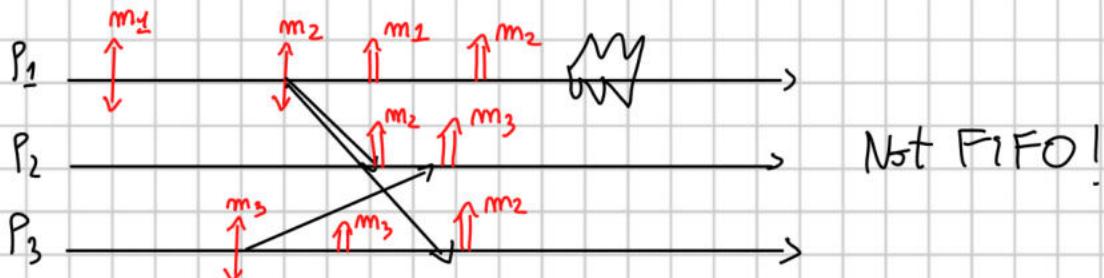
We have Perfect link, & perfect failure detector and
 in init: $mySm = 0$ and $\text{delivered} = \emptyset$

1. XBroadcast implements a Best Effort Broadcast because we satisfy No duplication, Validity, No creation given that we have PP2P link, but don't satisfy RB because a message sent by a faulty process could be delivered only by a part of correct processes, we don't retransmitt.

Example : \uparrow X-broadcast \uparrow X-Deliver



2. Considering only the ordering, we have no guarantee on the ordered because we do a simple broadcast, the sequence number is not used, given that don't satisfy nothing. This example is only BEB:



Exercise 4

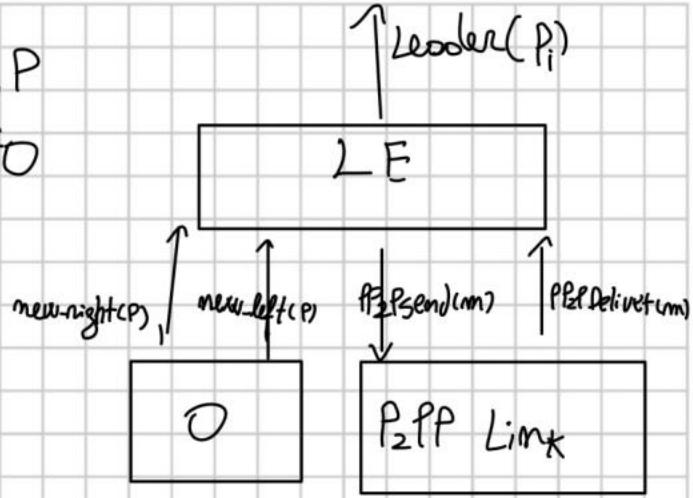
$\Pi = \{P_1, \dots, P_m\}$ M processes, FIFO PP2P link, we have a line:



P_i can communicate only with P_{i-1} and P_{i+1} . We have on each with primitive next-right (P) and next-left (P). Implement a perfect failure detector.

implements: Perfect Failure Detector, P

Uses: Oracle O, P2P link PL FIFO



Upon event $\langle P, \text{init} \rangle$ do

LEFT = get-left()

RIGHT = get-right()

alive = \top

detected = \emptyset

Upon event $\langle O, \text{new-right } | p \rangle$ do

Ex-right = RIGHT

RIGHT = P

detected = detected $\cup \{Ex\text{-RIGHT}\}$

trigger crash(Ex-right)

IF RIGHT \neq null do

trigger PP2PSend(|CRASH|, ExRight) to RIGHT

IF LEFT \neq null

trigger PP2PSend(|CRASH|, ExRight) to LEFT

Upon event $\langle O, \text{new-left } | p \rangle$ do

Ex-left = LEFT

LEFT = P

detected = detected $\cup \{Ex\text{-LEFT}\}$

trigger crash(Ex-left)

IF RIGHT \neq null do

trigger PP2PSend(|CRASH|, detected) to RIGHT

IF LEFT \neq null

trigger PP2PSend(|CRASH|, detected) to LEFT

Upon event $PP_2P\text{Deliver}(|CRASH|, P_i)$ from P_j

$\text{detected} = \text{detected} \cup P_i$

~~trigger~~ $\text{crash}(P_i)$

If $P_j = \text{LEFT}$ AND $P_j \neq \text{null}$

~~trigger~~ $PP_2P\text{Send}(|CRASH|, \text{detected})$ to RIGHT

else if $P_j = \text{RIGHT}$ AND $P_j \neq \text{null}$

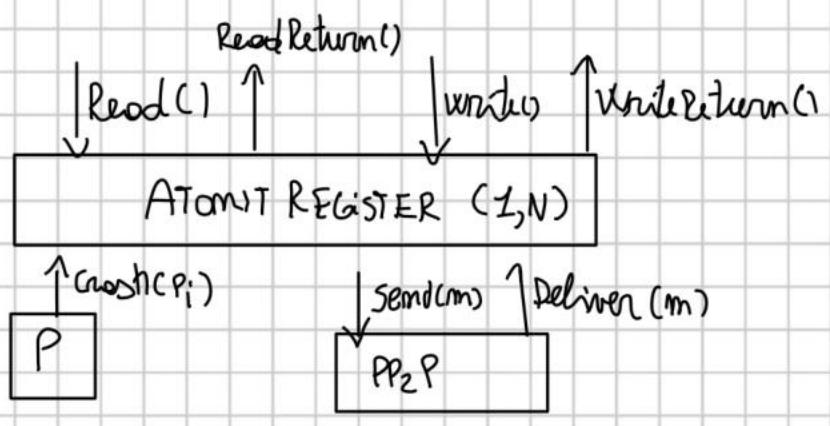
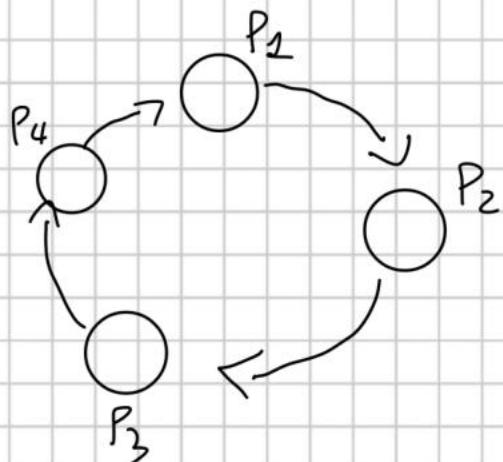
~~trigger~~ $PP_2P\text{Send}(|CRASH|, \text{detected})$ to LEFT

X NO MULTIPLE
CRASH ADD

In this way we satisfy Strong Completeness and Accuracy given that we have PP_2P link and the oracle.

Exercise 5

Ring topology with N processes. We have PP_2P link and a perfect failure detector P :



Idea is to write the value on a token and send it in the ring, invoke $\text{writeReturn}()$ only when turn back, and can read a process only when have the token. When a process crash and is our next, change next & recreate a new token.

Use sequence number for discard past token

implements: (1-N)-Atomic Register , AR

Users: PP2P link , PL

Perfect Failure Detector , P

int :

alive = π

val = \perp , sm = 0

sequence-number of token sem

Reading = False

Write = False

next = $P_{(i+1) \bmod N}$

If self = P_0

sm = sm + 1

TOKEN = val

trigger PP2PSend (<|READING|, sm>, TOKEN) to next

Upon event <AR, Write | v> do

write = TRUE

→ move when token can WRITE

val = v

and wait for return

Upon event <AR, Read> do

Reading = TRUE

Upon event PP2P Deliver (<|READING|, s>, TOKEN)

If $s = sm$

do nothing , have received same TOKEN

else if WRITE = TRUE

TOKEN = val , sm = sm + 1

trigger PP2PSend (<|WRITING|, sm>, TOKEN) to next

else if READING = TRUE

val = TOKEN , READING = FALSE , SM = SM + 1

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

trigger Read Return (val)

else

val = TOKEN, sm = sm + 1

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

Upon event PP2P Deliver (<|WRITING1, ss>, TOKEN)

if sm = s

do nothing, have received same TOKEN

else if write = TRUE

sm = sm + 1

trigger Write Return ()

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

else if READING = TRUE

val = TOKEN, READING = FALSE, sm = sm + 1

trigger PP2P Send (<|WRITING1, sm>, TOKEN) to next

trigger Read Return (val)

else

val = TOKEN, sm = sm + 1

trigger PP2P Send (<|WRITING1, sm>, TOKEN) to next

Upon event Clash (p_i)

alive = alive / {p_i}

if p_i = next

return the alive next

next = next_alive ()

trigger PP2P Send (<|READING1, sm>, TOKEN) to next



if next have same sm have received
the precedent token

Exercise 6

1.

No duplication: is not guarantee because if a process that have a set of delivery, it is affected by a transient failure and his variable FromEP is changed with other values, there is the possibility that in ϵ be delivered the condition $m \notin \text{From}[S]$ is satisfied multiple times and have multiple delivery of some message.

No creation: is not guarantee because if a process is affected by a transient failure, its FromEP is changed randomly, also with values not given by a broadcast. If the same process in the same interval time crash, then the other correct processes may deliver some wrong messages that are in FromEP, not previously broadcast by a process.

Validity: is guarantee, because we consider a transient failure processes as actually failed, if a correct process p broadcast a message m, thanks to " $\text{if } m \notin \text{From}[S]$ ", p will deliver it at some time.

Agreement: is guarantee because if we consider a transient failure processes as actually failed, all the correct processes will deliver the same set of messages, but some messages could be not satisfy no creation. This is guarantee by the From[S] variable that

check if the message is already delivered and by the fact that in every crash we broadcast all messages received from that processes, in this way all correct processes receive missing messages.

2. After t_{stab} the processes will not be affected by transient failure, but some processes are been effected, the algorithm is the same, now the precedent processes that was affected by transient failure return correct, for these also validity and agreement is not eventually guarantee.

Validity: if in the transient failure of p , $f_{from[p]}$ is changed with randomly messages and after t_{stab} that process want to broadcast a message and is already in $f_{from[p]}$, he will not deliver it, violating agreement.

Agreement: we have the same problem of validity, for example like we say for validity p could not deliver message m , but other correct processes could deliver it, and we have that correct processes have different set of messages, violating agreement.

3. For guarantee the No Duplication, validity and agreement after t_{stab} we can use the assumption that by N processes F_t can suffer transient failure. We use a Quorum

with $N - F_C$ (faulty not transient) messages, we send with every messages the $f_{\text{from}}[p]$ variable, when we receive $N - F_C$ from $[p]$ we do an interception between all $f_{\text{from}}[p]$, and eventually we achieve all the property, because the missing messages are retransmitted and the randomizing is correctly.

$$E[R] \stackrel{\text{MM1}}{=} \frac{1}{100-75} + \frac{1}{160-95} + \frac{1}{90-75}$$

Exercise 7

3 components handles sequentially the service.

$\lambda = 75 \text{ req/sec}$ is the arrival rate

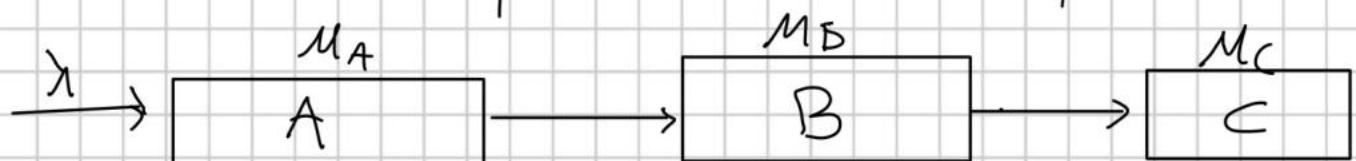
$$M_A = 100 \text{ req/sec}$$

$$M_B = 80 \text{ req/sec}$$

$$M_C = 90 \text{ req/sec}$$

1) double one component employing a perfect load balancer.

2) substitute one component with one 40% faster.



1) in the first case we have a perfect load balancer, we double the slowest component $M_B = 160 \text{ req/sec}$.

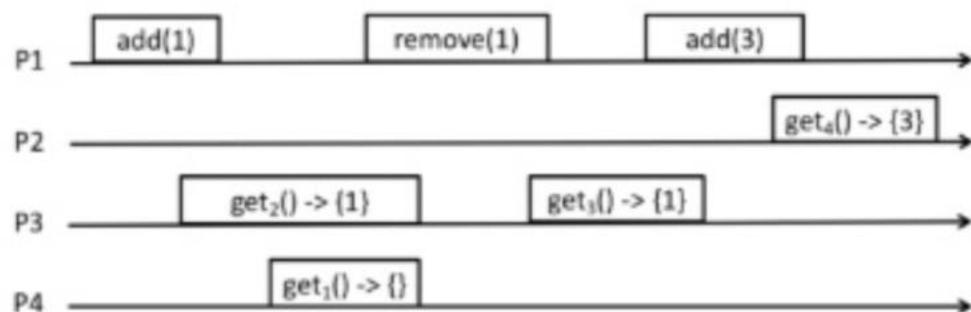
$$E[R] \stackrel{\text{MM1}}{=} \frac{1}{100-75} + \frac{1}{160-95} + \frac{1}{90-75} = \\ = 0,045 + 0,011 + 0,066 = 0,117 \text{ s}$$

2) With second option upgrade always the slowest component and obtain $M_B = 112 \text{ req/sec}$

$$E[R]^{M/M/1} = \frac{1}{25} + \frac{1}{37} + \frac{1}{15} = \\ = 0,045 + 0,0275 + 0,065 = 0,1375$$

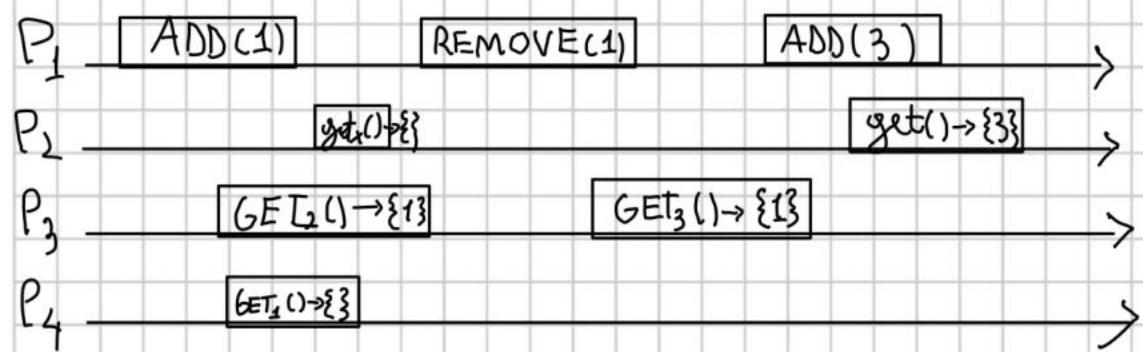
We chose option 1) and upgrade component B!

Exercise 8



1. No is not linearizable because of $\text{get}_1() \rightarrow \{\}$ and $\text{get}_3() \rightarrow \{3\}$ are in contradiction, we have no problem with $\text{get}_1() \rightarrow \{\}$, $\text{get}_2() \rightarrow \{1\}$ and $\text{remove}(1)$ because are concurrent and we can rearrange, but $\text{get}_3()$ we can drop only with $\text{remove}(1)$. we can not provide a sequence is legal.

2. The execution with the get-validity provided in the text, it is linearizable, we provide a new execution that is not linearizable and satisfy get validity.



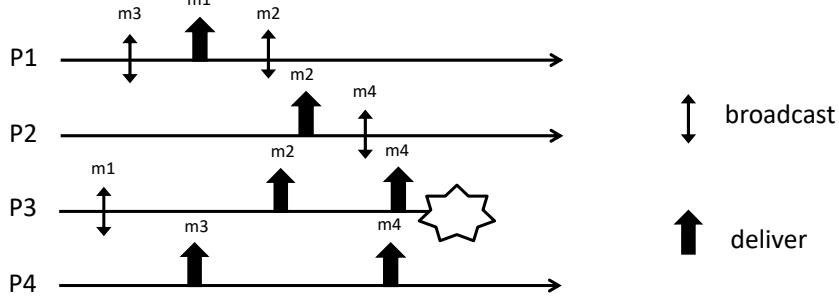
We change the execution, now we have three concurrent events get_x , get_2 and get_1 but get_x violate validity because return 1 but no after $ADD(1)$.

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 19 – Exercises
November 16th, 2022

Ex 1: Consider the partial execution depicted in the Figure

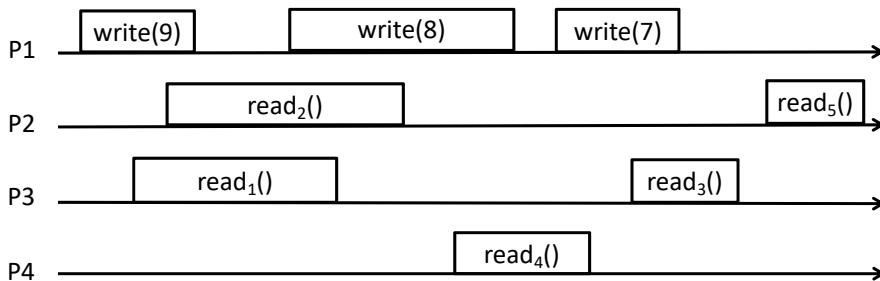


Answer to the following questions:

1. Provide ALL the possible delivery sequences that satisfies causal order and TO (UA, SUTO)
2. Complete the execution in order to have a run satisfying TO (UA WNUTO), FIFO order Broadcast but not Causal Order Broadcast
3. Complete the execution in order to have a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast and not satisfying Total Order.

NOTE: In order to solve the exercise, you can only add broadcast, deliveries and failures.

Ex 2: Consider the partial execution depicted in the Figure



Answer to the following questions:

1. Define ALL the values that can be returned by read operations (Rx) assuming that the run refers to a regular register.
2. Define ALL the values that can be returned by read operations (Rx) assuming that the run refers to an atomic register.

3. Assign to each read operations (R_x) a return value that makes the execution linearizable.

Ex 3: Let us consider the following algorithm

```

upon event < $frb$ , Init > do
   $lsn := 0;$ 
   $pending := \emptyset;$ 
   $next := [1]^N;$ 

upon event < $frb$ , Broadcast |  $m$  > do
  for each  $p \in \Pi$  do
    trigger < $Send$  | [DATA, self,  $m$ ,  $lsn$ ] > to  $p$ ;
   $lsn := lsn + 1;$ 

upon event < $Deliver$  |  $p$ , [DATA,  $s$ ,  $m$ ,  $sn$ ] > do
   $pending := pending \cup \{(s, m, sn)\};$ 

  while exists  $(s, m', sn') \in pending$  such that  $sn' = next[s]$  do
     $next[s] := next[s] + 1;$ 
     $pending := pending \setminus \{(s, m', sn')\};$ 
    trigger < $frb$ , Deliver |  $s, m'$  >;
  
```

Let us consider the following properties:

- **Validity:** If a correct process p broadcasts a message m , then p eventually delivers m .
- **No duplication:** No message is delivered more than once.
- **No creation:** If a process delivers a message m with sender s , then m was previously broadcast by process s .
- **Agreement:** If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- **FIFO delivery:** If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

Assuming that every process may fail by crash, address the following points:

1. Considering that messages are sent by using *perfect point to point links*, for each property mentioned, discuss if it satisfied or not and provide a motivation for your answer:
2. Considering that messages are sent by using *fair loss links*, for each property mentioned, discuss if it satisfied or not and provide a motivation for your answer.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links and are structured through a ring (i.e., each process p_i can exchange messages only with processes $p_{i+1 \text{ mod } n}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $new_next(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Uniform Reliable Broadcast communication primitive.

Ex 5: Let us consider the following algorithm implementing a (1, N) atomic register in synchronous system.

```

1  upon event { onar, Init } do
2    (ts, val) := (0, ⊥);
3    correct := Π;
4    writeset := ∅;
5    readval := ⊥;
6    reading := FALSE;

7  upon event { P, Crash | p } do
8    correct := correct \ {p};

9  upon event { onar, Read } do
10   reading := TRUE;
11   readval := val;
12   trigger { beb, Broadcast | [WRITE, ts, val] };

13 upon event { onar, Write | v } do
    trigger { beb, Broadcast | [WRITE, ts + 1, v] };

```

```

14 upon event { beb, Deliver | p, [WRITE, ts', v'] } do
15   if ts' > ts then
16     (ts, val) := (ts', v');
17     trigger { pl, Send | p, [ACK] };

18 upon event { pl, Deliver | p, [ACK] } then
19   writeset := writeset ∪ {p};

20 upon correct ⊆ writeset do
21   writeset := ∅;
22   if reading = TRUE then
23     reading := FALSE;
24     trigger { onar, ReadReturn | readval };
25   else
26     trigger { onar, WriteReturn };

```

Assuming that messages are sent by using perfect point-to-point links and that the broadcast is best effort answer the following questions:

1. Discuss what does it happen to every atomic register property (i.e., termination, validity and ordering) if the failure detector is eventually perfect and not perfect
2. Discuss what does it happen to every atomic register property (i.e., termination, validity and ordering) if we change line 12 with **trigger** { *beb*, Broadcast | [WRITE, ts+1, val] };

Ex 6: Consider a distributed system composed of n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links. Processes are connected through a directed ring (i.e., each process p_i can exchange messages only with processes $p_{i+1 \text{(mod } n)}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $\text{new_next}(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Leader Election primitive at every process p_i .

Exercise 1

1. $m_3 \rightarrow m_2 \rightarrow m_4$, $m_1 \rightarrow m_2$

$m_3 \rightarrow m_4$ FOT TO

m_3, m_1, m_2, m_4

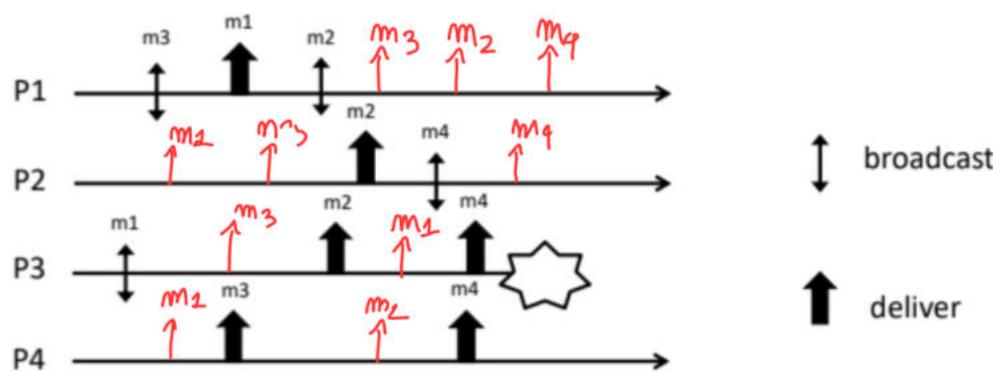
m_1, m_3, m_2, m_4

must be the possible sequences for all processes for satisfy TO(UA, SUT) and causal order.

2. $m_3 \rightarrow m_2$ FIFO

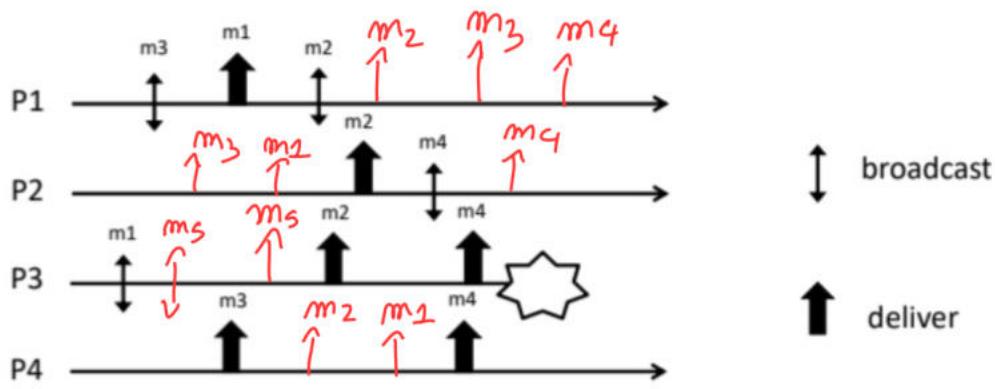
$m_1 \rightarrow m_2 \rightarrow m_4$ CAUSAL ORDER

$m_3 \rightarrow m_4$



Satisfy TO(UA, UNTO) and FIFO order broadcast but not causal, because all correct processes have some set of delivery and also faulty P3 this satisfies UA, UNTO because correct processes have some order of delivery but the faulty have an inversion ($m_3 \rightarrow m_1$). Not causal because in P3 we have $m_1 \rightarrow m_2$.

3. We show on execution that satisfy regular reliable broadcast but not uniform and not total order.



Without adding a broadcast in P₃ and the delivery of m₅ we can not have RB but not have UA, because we must have that the faulty process have a different set of messages respect to corrects one.

Exercise 2

1. Rx values with regular register

R₂: 0, 9, 8 , R₁: 0, 9, 8 , R₃: 8, 7 , R₄: 9, 8, 7 , R₅: 7

2. Rx values with an atomic register

R₁: 0, 9, 8 , R₂: 0, 9, 8 , R₅: 7

R₄: 9, 8, 7 or 8, 7 because if R₁ or R₂ return 8 then R₄ can not return 9.

R₃: 8, 7 or 7 because if R₄ return 7 then R₃ can not return 8 .

3. With these assignment this execution is linearizable:

R₁: 0 , R₂: 9 , R₄: 8 , R₃: 7 , R₅: 7

We respect causal order and also validity.

Exercise 3

1. Considering we have P2P link:

Validity: is not satisfied because all processes set $lm = 0$ and $next[1]^N$, when they do the first broadcast send the message with $sm = 0$ and increment lm only after the send. The message will be added in pending but never delivered because we never satisfy $0 = next[5]$.

No duplication: is satisfied because we use sequence number for messages, send the messages only one time in the broadcast and we continuously check if we have something in pending with $sm = next[5]$ and increment sm when deliver the message.

No creation: is satisfied because we have P2P link that guarantee this property and also by the algorithm we add in pending only messages sent by a previous broadcast.

Agreement: is not satisfied because if a faulty process broadcast a message m and is added in pending by some correct processes, they will deliver it, but they don't retransmit it or neither when a process crash, and for this if a correct process don't receive the same broadcast it never deliver m .

FIFO delivery: is not satisfied because the second message sent by any process with $sm = 1$ will be always delivered before the first because it will remain forever in pending with $sm = 0$, for the other

messages after the second the property is satisfied, but the second violate FIFO delivery.

2. Using fair loss links:

Validity: Some is before and also we don't have the guarantee that when we send the message it will arrive at destination, and the algorithm doesn't have any retransmission.

No duplication: is guarantee because we have the sequence number and the check in the while if we succeed in send the message in FLL. Otherwise we have no problem because we don't send the message and can not be delivered.

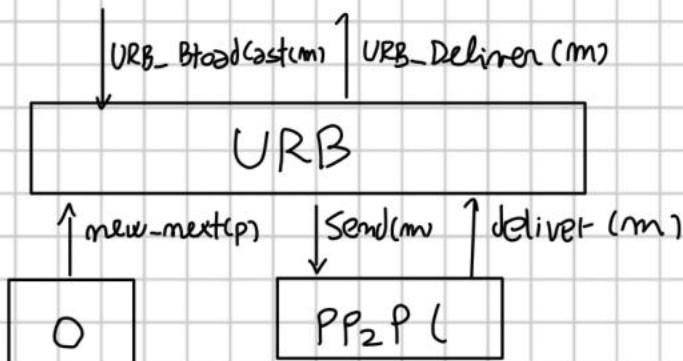
Agreement: is not satisfied like before for the problem with faulty processes and also we can lost the messages on the channel for these some faulty processes can deliver a message m but other correct processes could not deliver it.

No creation: is guaranteed by the FLL and by the algorithm itself

FIFO Delivery: is not satisfied because for the same problem that we show in 1., the second message is deliver before the first.

Exercise 4

m processes : $\Pi = \{P_1, \dots, P_m\}$, unique identifiers, we use P2P links and we have a ring structure. We can have crashes. we have an oracle with $\text{next_next}(p)$ when neighbor is failing. Implementing leader election for each P_i .



We use a token, every process that want to do a broadcast must have the token, and can deliver only when the token has complete one loop. We manage the crash retransmitting and check if the TOKEN is a duplication using the last delivery tag and SM.

Implementation: Uniform Reliable Broadcast, URB

Used: Perfect Point-to-point link, PP2P link

Oracle, O

init:

$\text{pending} = \perp$

$\text{next} = P_{i+1 \pmod m}$

$\text{delivered} = \emptyset$

$\text{wait} = \text{FALSE}$

$\text{S.M.} = 1 \longrightarrow \text{sequence number}$

$\text{last_delivered} = \perp$

If $\text{self} = P_1$ do

$\text{TOKEN} = \perp$

{Trigger $\text{PP2PSend}(\text{FREE } (\text{TOKEN}, \text{SM}))$ to next}

Upon event URB-Broadcast(m)

WAIT = TRUE → wait for token for Broadcast M

Pending = m

Upon event P2P Deliver (FREE1, <TOKEN, S>)

Sm = S + 1

IF WAIT = TRUE do

TOKEN = Pending,

trigger PP2PSend (BROADCAST1, <TOKEN, Sm>) to next

Pending = 1

else

trigger PP2PSend (FREE1, <TOKEN, S>) to next

Upon event P2P Deliver (BROADCAST1, <TOKEN, S>)

IF TOKEN = last-delivered do

IF Sm > S

< DISCARD> TOKEN is a DUPLICATIONS

else

Sm = S + 1

trigger PP2PSend (FREE1, <TOKEN, S>) to next

else IF WAIT = TRUE and pending = 1

WAIT = FALSE

delivered = delivered ∪ TOKEN

last-delivered = TOKEN

trigger URB-Deliver (TOKEN)

TOKEN = 1, Sm = S + 1

trigger PP2PSend (FREE1, <TOKEN, Sm>) to next

else

$$sm \leftarrow s + 1$$

$$\text{delivered} = \text{delivered} \cup \text{TOKEN}$$

$$\text{last-delivered} = \text{TOKEN}$$

trigger URB-DELIVR(TOKEN)

trigger PP2PSend(|BROADCAST|, <TOKEN, sm>) to next

Upon event next-next(p)

$$next = p$$

$$\text{TOKEN} = \text{last-delivered}$$

trigger PP2PSend(|BROADCAST|, <TOKEN, sm>) to next

We guarantee no duplication of TOKEN
checking CAST-DELIVER and the sequence number,
if the next that received the new TOKEN finds
that his sm is greater than the sender, discard
the TOKEN.

Exercise 5

(1-N) atomic register, P2P Links and BEB.

1. we have eventually failure detector:

Termination: is satisfied because if a correct process is suspected for error then the condition $\text{correct} \subseteq \text{writerset}$ is satisfied with less ACK, but the write or read will reach the termination.

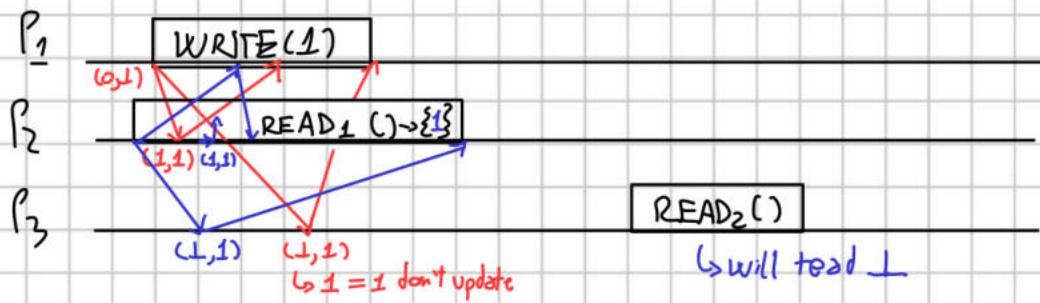
Validity: is not satisfied because if a process p that want to read is suspected for error and a process obs a `WriteReturn` because believe that p is crashed and also p doesn't receive the new value because writer has used BEB, it is possible that p that is not concurrent with last write, it read the precedent value violating validity.

Ordering: is not satisfied for the same motivation of validity. A correct process p that is been suspected would not receive the other messages and if he read a value could return a value that violating the correct order.

2. in line 12: `trigger <Beb, Broadcast [WRITE, ts+1, val]`

Termination: is satisfied because we change only the timestamp, the process will receive the ACK like before the change and satisfied " $\text{correct} \subseteq \text{writerset}$ "

Validity: is not satisfied because if we have this situation:



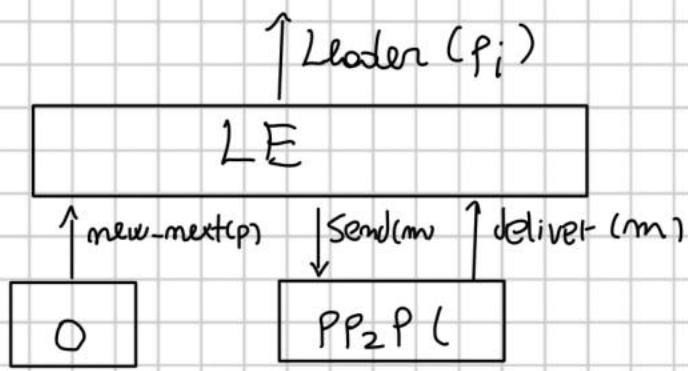
We can have 2 process after a write that return a value written before it, violate validity.

Agreement: is not satisfied like we see in the example
we can have a read_2 that is after read_1 that return a value that violate the ordering.

Exercise 6

m processes : $\Pi = \{P_1, \dots, P_m\}$, unique identifiers, we use P2P links and we have a ring structure. we can have crashes. we have an oracle with $\text{new-meet}(p)$ when neighbor is failing.

Implementing Uniform Reliable Broadcast:



We initialize P_m as the leader for all processes. When a process receives $\text{new-meet}(p)$, check if it is the leader, the next that is crashed, then if it is true become himself the leader and send to other the communication. otherwise send the actual leader to next.

then if it is true become himself the leader and send to other the communication. otherwise send the actual leader to next.

Implementation: Leader Election, LE

Used: Perfect Point to point link, P2P link
Oracle, O

init:

$$\text{next} = P_{(i+1) \bmod n}$$

$$\text{alive} = \Pi$$

$$\text{leader} = P_m$$

$$\text{WAIT} = \text{FALSE}$$

Upon event new_next (p)

IF leader > p AND self > p

leader = self

trigger leader (leader)

next = p

Wait = TRUE \rightarrow Wait all confirm new leader

trigger PP2PSend (NEW-LEADER, leader) to next

else

next = p

trigger PP2PSend (NEW-LEADER, leader)

Upon event PP2PDeliver (NEW-LEADER,)

IF leader = ()

IF wait = TRUE

wait = FALSE

< Stop resending m, i'm new leader >

else

< Discard message, it is already received the messages >

else

leader = ()

trigger leader (leader)

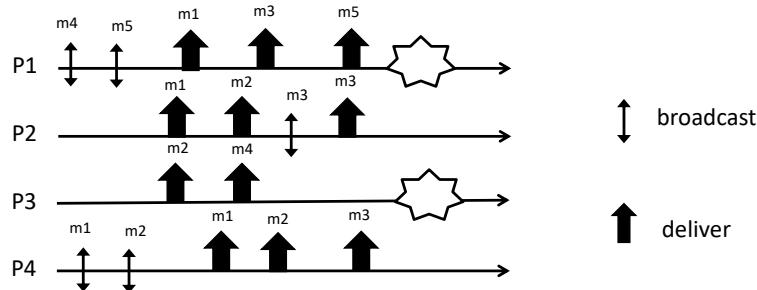
trigger PP2PSend (NEW-LEADER, leader) to next

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 23 – Exercises
November 23th, 2022
(Estimated time to complete all exercises: 3 hours)

Ex 1: Consider the execution depicted in the Figure

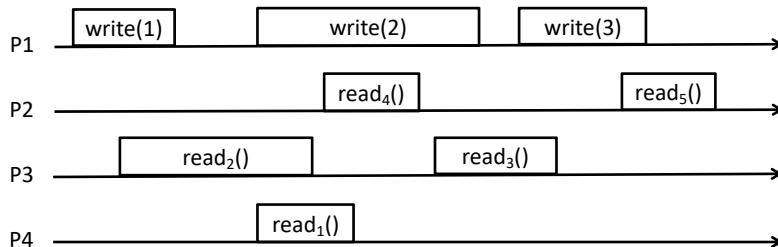


Answer to the following questions:

1. Which is the strongest TO specification satisfied by the proposed run?
Motivate your answer.
2. Does the proposed execution satisfy Causal order Broadcast, FIFO Order Broadcast or none of them?
3. Modify the execution in order to satisfy TO(UA, WUTO) but not TO(UA, SUTO).
4. Modify the execution in order to satisfy TO(NUA, WNUTO) but not TO(UA, WNUTO).

NOTE: In order to solve point 3 and point 4 you can only add messages and/or failures.

Ex 2: Consider the execution depicted in the following figure and answer the questions



1. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.

2. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.
3. Let us assume that values retuned by read operations are as follow: $\text{read}_1() \rightarrow 2$, $\text{read}_2() \rightarrow 2$, $\text{read}_3() \rightarrow 3$, $\text{read}_4() \rightarrow 1$, $\text{read}_5() \rightarrow 3$. Is the run depicted in the Figure linearizable?

Ex 3: Consider the algorithm shown in the Figure

<pre> upon event { Init } do delivered := Ø; pending := Ø; correct := Π; forall m do ack[m] := Ø; upon event { urb, Broadcast m } do pending := pending ∪ {(self, m)}; trigger { beb, Broadcast [DATA, self, m] }; upon event { beb, Deliver p, [DATA, s, m] } do ack[m] := ack[m] ∪ {p}; if (s, m) ∈ pending then pending := pending ∪ {(s, m)}; trigger { beb, Broadcast [DATA, s, m] }; </pre>	<pre> upon event {◊P,Suspect p} do correct := correct \ {p}; upon event {◊P,Restore p} do correct := correct ∪ {p}; function candeliver(m) returns Boolean is return (correct ⊆ ack[m]); upon exists (s, m) ∈ pending such that candeliver(m) do delivered := delivered ∪ {m}; trigger { urb, Deliver s, m }; </pre>
---	--

Assuming that the algorithm is using a Best Effort Broadcast primitive and an Eventually Perfect Failure Detector $\diamond P$ discuss if the following properties are satisfied or not and motivate your answer

- *Validity*: If a correct process p broadcasts a message m, then p eventually delivers m.
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s, then m was previously broadcast by process s.
- *Uniform agreement*: If a message m is delivered by some process (whether correct or faulty), then m is eventually delivered by every correct process.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_N\}$ with unique identifiers that exchange messages through perfect point-to-point links and are structured in a ring topology (i.e., each process p_i can exchange messages only with processes $p_{(i+1)\bmod N}$ and stores its identifier in a local variable `next`).

Each process p_i knows the initial number of processes in the system (i.e., every process p_i knows the value of N).

1. Assuming that processes are not going to fail, write the pseudo-code of an algorithm that implements a (1, N) regular register.

Ex.5: Answer true or false to the following claims providing a motivation:

1. The performance of a system can be analyzed independently from its load.
2. Let us assume a single service with a single class of requests (i.e. a single workload component). If we are under the stability condition ($\lambda < \mu$) then the expected response time of the system is independent from the arrival pattern.
3. The workload parameters that mostly influence the performance of system are the arrival pattern and the service demands.
4. The time needed to run a simulation is upper-bounded by the simulated time
5. The reliability function $R(t)$ associated to a component may increase after the restoration of a component

Ex.6: A service is provided through 3 types of components (e.g. database, webserver, and an application), referred with A, B and C. At least one instance of each component must be available to provide the service. One only instance is available for components A and C, whereas two instances of component B are available.

Assuming that all failures and restorations are independent among them, that the failures rates of the three components A, B and C are respectively 0.2 , 0.5 and 0.3 faults per day, that the mean time to repair the components A, B, and C are respectively 2h, 3h, and 1.5h, is the service available at least 98% of the time? If not, is the availability target met deploying an additional instance of one of the components?

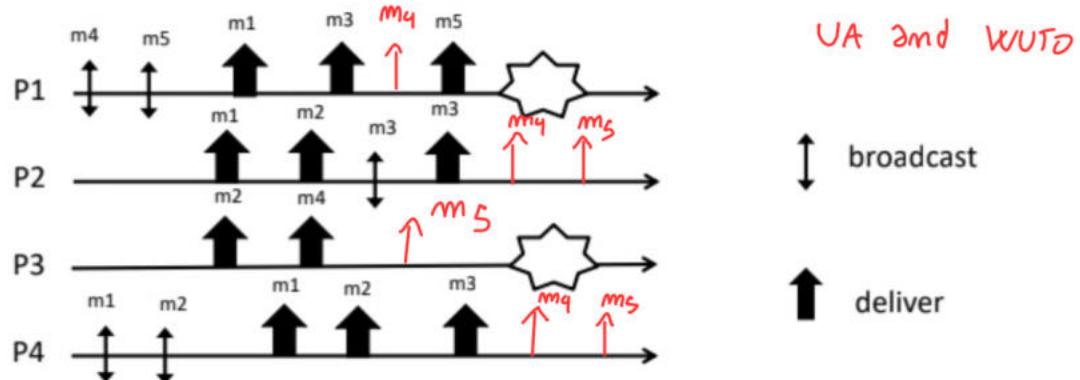
Ex.7: Let us assume a probabilistic version of the Eager Reliable Broadcast protocol in which every process retransmits a rb-delivered message only once with probability 0.5. The rb-broadcast operation is periodically triggered, and the calls follows an exponential distribution with average time between consecutive calls of 10 seconds. Assuming a distributed system of 30 processes and no processing delay, what is the minimum in-channel capacity per process (in terms of message per seconds) to ensure that the expected time for a rb-delivery is below 0.5 seconds?

Exercise:

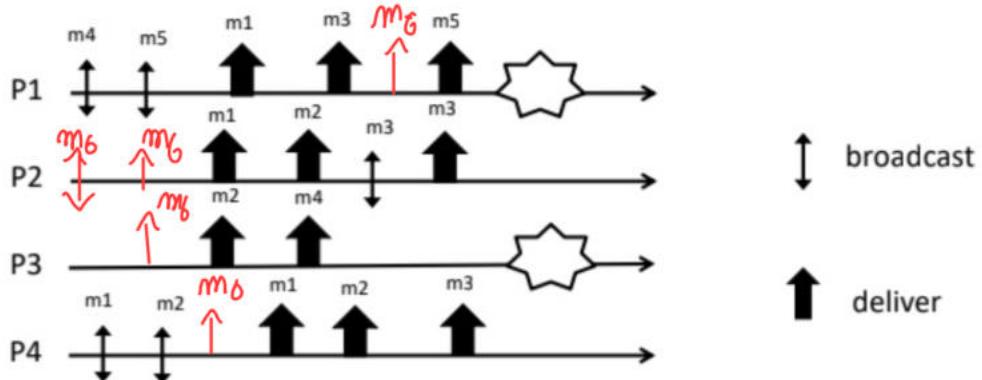
1. We don't have UA because P_1 is faulty process have a set of delivery that is not a subset of the correct one, for m_5 . We have WUTO because the order of pair is respected but not SUTO because we have holes in delivery:
TO(NUA, WUTO)

2. For Fifo we have $m_4 \rightarrow m_5$, $m_1 \rightarrow m_2$
For causal order also $m_1 \rightarrow m_3$ and $m_2 \rightarrow m_3$
Causal order and FIFO broadcast are satisfied because the order in delivery is satisfied in all correct processes.

3.



4.



Exercise 2

1. $R_1 = \{1, 2\}$, $R_2 = \{0, 1, 2\}$, $R_3 = \{1, 2, 3\}$, $R_4 = \{1, 2\}$, $R_5 = \{2, 3\}$
2. $R_2 = \{0, 1, 2\}$, R_4 depend on R_2 if $R_2 = 2$ then $R_4 = 2$ else $R_4 = \{1, 2\}$, R_1 is concurrent with R_2 and R_4 then can return 1, 2 independently from the other 2, R_3 depend on the value before if R_1 or R_4 or R_2 return 2 then $R_3 = \{2, 3\}$ else $\{1, 2, 3\}$, R_5 depend on R_3 if $R_3 = \{3\}$ then $R_5 = \{3\}$ else $\{2, 3\}$.
3. $R_1 = 2$, $R_2 = 2$, $R_3 = 3$, $R_4 = 1$, $R_5 = 3$ is not linearizable because R_2 and R_4 are not concurrent R_4 return 1 after $R_2 = 2$, violating linearizability.

Exercise 3

Validity: is satisfied because when we broadcast a message p we add it in pending, and wait that the ACK associated to that message is equal to the number of correct processes. We have two case we wrongly suspect a wrong process and deliver the message or we receive all ACKs and deliver, but in any case we deliver p .

No duplication: is not guaranteed because we use pending and don't delete any messages from it, if we wrongly suspect a process, we would deliver the message without an ACK, but when restore the process and arrive his ACK the undeliver(p) check is true, and given that the message is in pending, we will deliver it again.

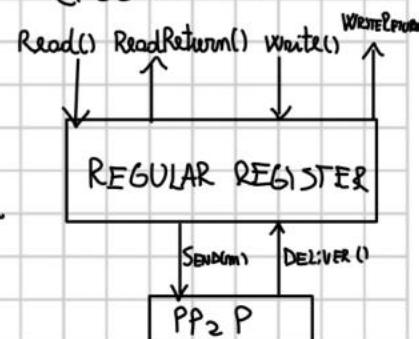
No Cetision: also is guarantee we deliver only messages in pending that by the algorithm have only messages previously broadcasted.

Uniform agreement: is violated because if for example we wrongly suspect all other processes and in that period we broadcast a message p , we deliver after receive the delivery of Beb , but when we restore the other correct processes we don't retransmit again p then we have a delivery that other correct processes don't have.

Exercise 4

n processes $\Pi = \{p_1, p_2, \dots, p_n\}$, P2P link, each p_i have next $= p_{(i+1) \bmod n}$, No failure.

We have only a writer for example p_1 that send the value and wait for return for do the writeReturn, the other can read their local value at any time.



Implementation: Regular Register ($1, n$), RR

Used: Perfect Point-to-Point link, P2P Link

init:

$val = \perp$

$next = p_{(i+1) \bmod n}$

Upon event $\langle RR, Read \rangle$ do

trigger $\langle RR, ReadReturn | val \rangle$

upon event <RR, write, v> do

val = v

trigger PP2PSend (IWRITEI, val) to next

upon event PP2PDeliver (IWRITEI, v) from next

IF val >= v do

trigger writeReturn <val>

else

val = v

trigger PP2PSend (IWRITEI, val) to next

When we want read, we read local value and do readReturn. The writer update value, and send the value in the ring, when return the value to the write return, each reader simply update value and send to next.

Exercise 5

1. False, the performance of a system depend heavily on the characteristics of its load, because the expected performances depend on the characteristic of the inputs, we must define very well the workload.

2. λ is not enough to characterize the workload of a system, the response time depend on the arrival pattern \Rightarrow False

3. True, they are fundamental, because they characterize the arrival pattern, given the information about queuing. The service demands describe the demand on the system.

The service time per completion is important but it is part of service demands, and the response time is not part of workload

4. False, generally there is no relationship between simulated time and the time needed to run a simulation on the computer, the simulation clock is different from the clock of computer.

5. False, Reliability doesn't depend on recovery.

Exercise 6

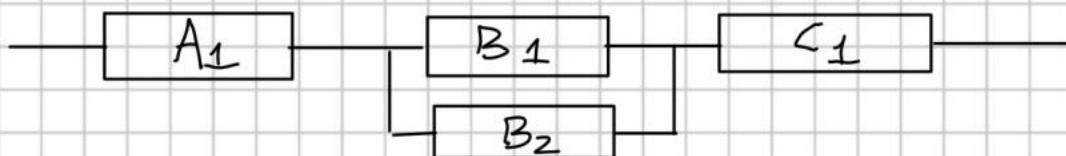
instance for each

3 components A, B, C. At least one available for provide the services. A, C one instance, B two instances.

$$\lambda_A = 0.2, \lambda_B = 0.5, \lambda_C = 0.3, MTTR_A = 2 \text{ h}, MTTR_B = 3 \text{ h}, MTTR_C = 1.5 \text{ h}$$

Service is available at least 98% of the time? we achieve it adding one instance to A or B or C?

Because service is guaranteed also if only one instance is available we have this situation:



$$A_{\text{serial}} = \prod A_i = A_A \cdot A_B \cdot A_C$$

$$A_{\text{parallel}} = 1 - \prod (1 - A_i) \Rightarrow A_B = 1 - (1 - A_{B1})(1 - A_{B2})$$

$$\lambda = \frac{1}{MTBF} \Rightarrow MTBF = \frac{1}{\lambda}, A_i = \frac{MTBF}{MTBF + MTTR}$$

$$MTBF_A = \frac{1}{0.2 \text{ ft/day}} = 120 \text{ h} \quad MTBF_B = \frac{1}{0.5 \text{ ft/day}} = 20 \text{ h}$$

$$MTBF_C = \frac{1}{0.3 \text{ ft/day}} = 33.3 \text{ h}$$

$$A_A = \frac{120}{2+120} = 0,983$$

$$A_{B_1} = A_{B_2} = \frac{40}{40+3} = 0,94$$

$$A_B = 1 - (1 - 0,94)^2 = 0,9964$$

$$A_C = \frac{80}{1,5+80} = 0,981$$

$$A_{TOT} = A_A \cdot A_B \cdot A_C = 0,96 < 0,98$$

Not with this inputs not satisfied 98%

Try to optimize the bottleneck C with two instances:

$$A_C = 1 - (1 - 0,981)^2 = 0,999$$

$$A_{TOT} = A_A \cdot A_B \cdot A_{C_{new}} = 0,98$$

We achieve objective adding a instance to C.

Exercise 7

Probabilistic Eager Reliable Broadcast, 1b-delivered retransmitted once with probability 0.5. 1b-broadcast is triggered every 10 seconds, 30 processes and no delay. message per second capacity per process what is the minimum capacity for it to ensure Expected time for a rb-delivery is less than 0.5 seconds?

We use a queue model with:

$$R = \text{EXPECTED TIME TO DELIVER} = \frac{1}{M_i - \frac{30}{\downarrow \text{PROCESSES}} \cdot \frac{1}{10} \cdot \frac{1}{2}} \xrightarrow{\text{L} \rightarrow \text{PROBABILITIES}}$$

$$\frac{1}{M_i - \frac{3}{2}} < \frac{1}{2} \Rightarrow \frac{2}{2M-3} < \frac{1}{2}$$

$$2M-3 > 4 \Rightarrow M > \frac{7}{2} = 3.5 \text{ messages/sec}$$

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 27 – Exercises
November 30th, 2022

Ex 1: Let us consider a distributed system composed of N processes executing the algorithm reported in figure

Algorithm 3.12: Broadcast with Sequence Number

Implements:

FIFOReliableBroadcast, **instance** *frb*.

Uses:

ReliableBroadcast, **instance** *rb*.

```
upon event <frb, Init > do
    lsn := 0;
    pending := {};
    next := [1]N;
    
upon event <frb, Broadcast | m > do
    lsn := lsn + 1;
    trigger <rb, Broadcast | [DATA, self, m, lsn] >;
    
upon event <rb, Deliver | p, [DATA, s, m, sn] > do
    pending := pending ∪ {(s, m, sn)};
    while exists (s, m', sn') ∈ pending such that sn' = next[s] do
        next[s] := next[s] + 1;
        pending := pending \ {(s, m', sn')};
        trigger <frb, Deliver | s, m' >;
```

Let us assume that

1. up to f processes may be Byzantine faulty and
2. A Byzantine process is not able to compromise the underline Reliable Broadcast primitive (i.e., when a Byzantine process sends a message through the *rbBroadcast* interface, the message will be reliably delivered to every correct process).

For each of the following properties, discuss if it can be guaranteed when $f=1$ and motivate your answer (also by using examples)

- *Validity:* If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication:* No message is delivered more than once.
- *No creation:* If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement:* If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- *FIFO delivery:* If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

Ex 2: Let us consider a distributed system composed of N processes executing the algorithm reported in figure

Algorithm 3.2: Lazy Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb ;
PerfectFailureDetector, **instance** \mathcal{P} .

```

upon event  $\langle rb, Init \rangle$  do
     $correct := \Pi$ ;
     $from[p] := [\emptyset]^N$ ;

upon event  $\langle rb, Broadcast \mid m \rangle$  do
    trigger  $\langle beb, Broadcast \mid [DATA, self, m] \rangle$ ;

upon event  $\langle beb, Deliver \mid p, [DATA, s, m] \rangle$  do
    if  $m \notin from[s]$  then
        trigger  $\langle rb, Deliver \mid s, m \rangle$ ;
         $from[s] := from[s] \cup \{m\}$ ;
        if  $s \notin correct$  then
            trigger  $\langle beb, Broadcast \mid [DATA, s, m] \rangle$ ;

upon event  $\langle \mathcal{P}, Crash \mid p \rangle$  do
     $correct := correct \setminus \{p\}$ ;
    forall  $m \in from[p]$  do
        trigger  $\langle beb, Broadcast \mid [DATA, p, m] \rangle$ ;

```

Let us assume that up to f processes may be Byzantine faulty with a symmetric behaviour i.e., a Byzantine process can change the content of the message it is going to send, but it cannot send different values to different processes when invoking the $bebBroascast$ (they can lie but in a consistent way).

For each of the following properties, discuss if it can be guaranteed when $f=1$ and motivate your answer (also by using examples)

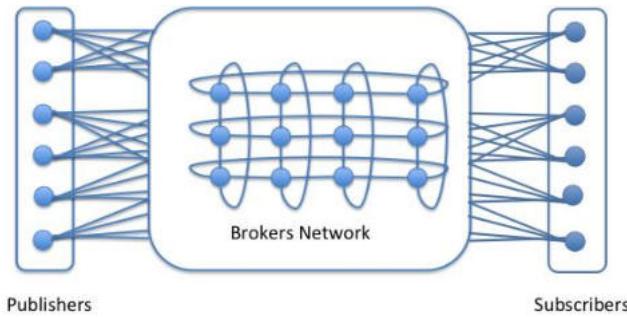
- *Validity:* If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication:* No message is delivered more than once.
- *No creation:* If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement:* If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Ex 3: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links and are arranged in a unidirectional ring (i.e., each process p_i can exchange messages only with process $p_{(i+1) \bmod n}$ and stores its identifier in a local variable `next`).

Each process p_i knows the initial number of processes in the system (i.e., every process p_i knows the value of n).

1. Assuming that processes are not going to fail, write the pseudo-code of an algorithm that implements a consensus primitive
2. Let's now assume that processes may crash and that each correct process has access to a perfect failure detector. Discuss the issues of the implementation provided in point 1 and describe how you should modify the algorithm to make it fault tolerant
3. Considering the algorithm proposed in point 1, discuss which properties may be compromised by the presence of one Byzantine process in the ring.

Ex 4: Let us consider a distributed system composed by publishers, subscribers and brokers. Processes are arranged in a network made as follows and depicted below:



1. Each publisher is connected to k brokers through perfect point-to-point links;
2. Each subscriber is connected to k brokers through perfect point-to-point links;
3. Each broker is connected to k brokers through perfect point-to-point links and the resulting broker network is k -connected (4 -connected in the example);

Answer to the following questions:

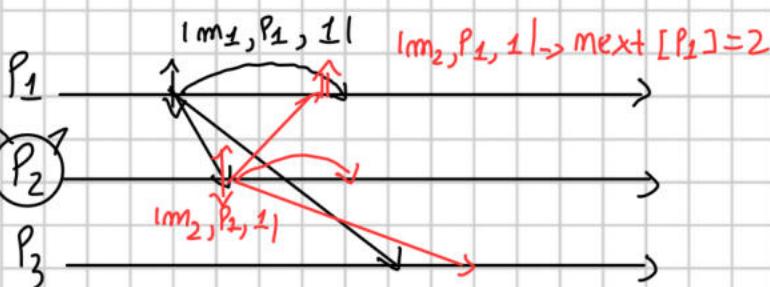
1. Write the pseudo-code of an algorithm (for publisher, subscribers and broker nodes) implementing the subscription-flooding dissemination scheme assuming that processes are not going to fail.
2. Discuss how many crash failures the proposed algorithm can tolerate.
3. Modify the proposed algorithm in order to tolerate f Byzantine processes in the broker network and discuss the relation between f and k .

Exercise 1

N processes, up to F processes may be Byzantine faulty and Byzantine process is not able to compromise the RB primitive. LINK ARE NOT AUTHENTICATED

For $F=1$:

Validity: is not guaranteed, because if a correct process P_1 send a broadcast a message m , $sm=1$, but the Byzantine process send a message m_2 with the same sm , like this:

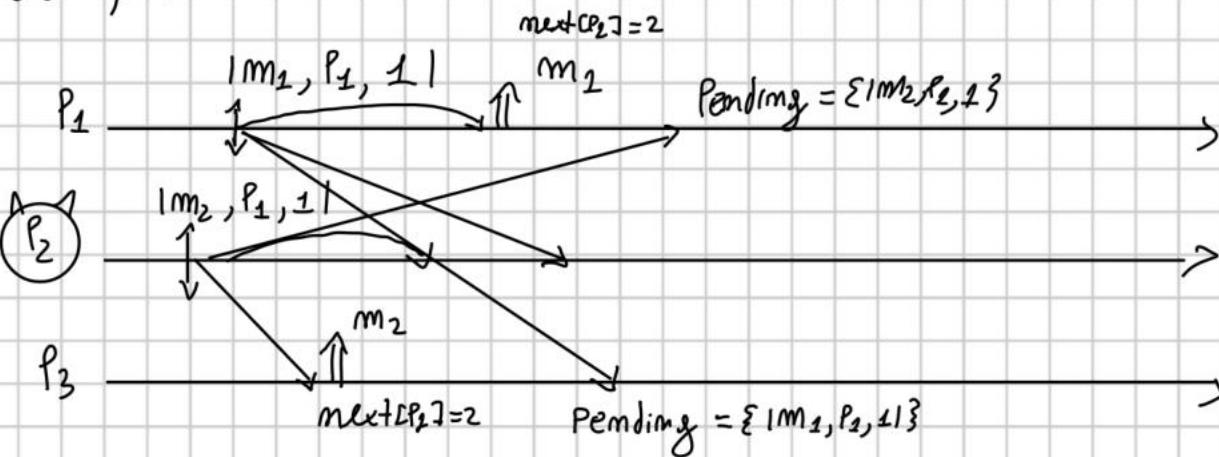


We have that P_2 receive $(m_2, P_2, 1)$, put it in pending and because $sm=1 = next[P_2]$, he will deliver it, but when receive $(m_1, P_1, 1)$ the original message will not be delivered because now $next[P_2] = 1$.

No duplication: is not guaranteed because like before a process P_1 broadcast a message $(m_1, P_1, 1)$, when the Byzantine process receive it, he can broadcast the message $(m_2, P_1, 2)$ that will be delivered by all correct processes because have the next sequence number, also the original message will be delivered with $sm=1$.

No creation: not guaranteed, the motivation is exactly the same for validity, Byzantine process can impersonate another process and achieve the delivery of the message.

Agreement: is not guaranteed because we can show this example:



P_1 deliver m_1 but P_3 deliver m_2 , violating that must deliver the same set of messages. The messages in pending will not be delivered because for the $SM=1$.

FIFO delivery: is not satisfied because byzantine process can broadcast a message m_3 with the same sequence number of m_1 , achieve the delivery and after m_2 will be delivered without m_1 .

Exercise 2

For $f = 1$:

Validity: it is satisfied because we have from $[p]$ variable and when we receive a deliver of a broadcast, we check that m is in the list, the byzantine process can only send the same message impersonating the correct process p or another message, but m is always delivered when we check $m \notin [s]$.

No duplication: is not guaranteed because if a process i deliver a message m , if the

Byzantine process send the same message but changing the sender, the process i deliver again the message because not from [S] not contain m with that sender.

No creation: not guarantee because we don't use authenticated links, the real sender p is not checked, the byzantine process can create false broadcast with fake sender.

Agreement: is guarantee because is true that byzantine can change the content of the message that he send, but have a symmetric behaviour, we achieve agreement thanks to the retransmission of messages of a crashed process.

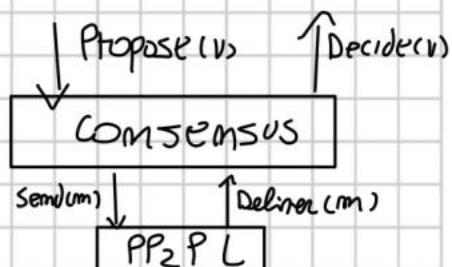
Exercise 3

n processes $\Pi = \{p_1, \dots, p_n\}$, P2P links, ring.

1. no failure, implements consensus primitive, we use a token where each process propose a value and after token is full we take the greater value.

Implementation: Consensus, C

Uses: P2P links, PC



init

state = wait

next = $p_{(i+1) \bmod n}$

$\Pi = n$

IF $\text{self} = p_0$:

$\text{TOKEN} = \text{TOKEN} \cup \{v\}$

trigger $\text{PP}_2\text{P}\text{Send}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{P}\text{Deliver}(\text{TOKEN})$

IF $\text{size}(\text{TOKEN}) == n$

$\text{decision} = \text{max}(\text{TOKEN})$

trigger $\text{Decide}(\text{decision})$

trigger $\text{PP}_2\text{P}\text{Send}(\text{decision})$ to next

else

$\text{TOKEN} = \text{TOKEN} \cup \{v\}$

trigger $\text{PP}_2\text{P}\text{Send}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{P}\text{Deliver}(\text{decision})$

IF $\text{State} == \text{WAIT}$

$\text{State} = \text{done}$

trigger $\text{Decision}(\text{decision})$

trigger $\text{PP}_2\text{P}\text{Send}(\text{decision})$ to next

else

<reach consensus>

2. We have a problem with crashes, messages can be lost and not reach consensus, we must add:

init

$\text{pending} = \emptyset$

$\text{decision} = \text{null}$

state = wait

next = $P_{(i+1) \bmod N}$

IT = n

IF self = P_0 :

TOKEN = TOKEN $\cup \{V\} \rightarrow \text{pend}_1 =$

pending = TOKEN my

trigger PP₂PSEND(TOKEN) to next

Upon event PP₂PDELIVER(TOKEN)

IF TOKEN == pending:

< discard duplicate token >

object

else IF size(TOKEN) == IT

number of c

decision = max(TOKEN)

processes

trigger Decide(decision)

trigger PP₂PSEND(decision) to next

else

TOKEN = TOKEN $\cup \{V\}$

Pending = TOKEN

trigger PP₂PSEND(TOKEN) to next

Upon event PP₂PDELIVER(d)

IF State == WAIT AND decision == null

State = done, decision = d

trigger Decision(d)

trigger PP₂PSEND(d) to next

else

< reach consensus > or < duplicate token >

Upon event crash(p)

$$\Pi = \Pi - 1$$

If $p == \text{next}$

$\text{next} = \text{next correct()}$

If $\text{decision} = \text{null}$

$\text{TOKEN} = \text{pending}$

trigger $\text{pp}_2 \text{Psend}(\text{TOKEN})$ to next

else

trigger $\text{pp}_2 \text{Psend}(\text{decision})$ to next

We have odd pending where we store the token value we have and decision where we store the decided value. In case of crash we update alive processes and if the next is crashed, we update next and resend the token or the decision. We check in deliver that the message is a duplicate.

3 Takking algorithm 1., we loose termination because the byzantine process can decide to not resend the token or the decision. We loose validity because byzantine process can change the TOKEN arbitrary and a process p can decide a value V never proposed. Integrity is guaranteed because correct processes have variable wait, we must add that is not uniform because of Byzantine. Agreement is not guaranteed because

if a process decide a value v and the next is the buy-in time process, it can send another decision V_1 .

Exercise 4 publish, subscribe, notify, unsubscribe

1. In subscription-flooding dissemination scheme, we have that each publisher when want to publish a value, he broadcast the value to all the K brokers which he is connected. The subscribers when want to subscribe send the message to each of the K broker which he is connected. When a broker receive a publish event start a notify event flooding and save the value, but before check if has saved a subscription with that value, and in affirmative we send the value to the subscriber. Subscriber get the notify event, and save the value.

Implementation: subscription-flooding

Uses: P2P links, PL

init:

delivered = \emptyset Publisher, subscriber and broker.

neighbours = $\{b_1, \dots, b_K\}$ Publisher and subscriber

↳ brokers can have links also with publisher and subscribers

subscription = \emptyset subscriber and broker

(subscription is an array of couples that are composed like (constraint, subscriber) where constraint is a function that take a value and return a boolean)

Upon event publish (v)

if $v \notin$ delivered :

delivered = delivered $\cup \{v\}$

for each $m \in$ neighbours : // send to blocks

trigger PP2PSend (NOTIFY, v) to m

Upon event subscribe (constraint)

subscription = subscription $\cup \{(\text{constraint}, \text{self})\}$

for each $m \in$ neighbours : // send to blocks

trigger PP2PSend (SUBSCRIPTION | <constraint, self>) to m

upon event PP2PDeliver (SUBSCRIPTION | <C, p>) from S

subscription = subscription $\cup \{(\mathcal{C}, p)\}$

↳ block save the subscription !

Upon event unsubscribe (constraint)

subscription = subscription / $\{(\text{constraint}, \text{self})\}$

for each $m \in$ neighbours :

trigger PP2PSend (UNSUBSCRIPTION | <constraint, self>) to m

upon event PP2PDeliver (UNSUBSCRIPTION | <C, p>) from S

subscription = subscription / $\{(\mathcal{C}, p)\}$

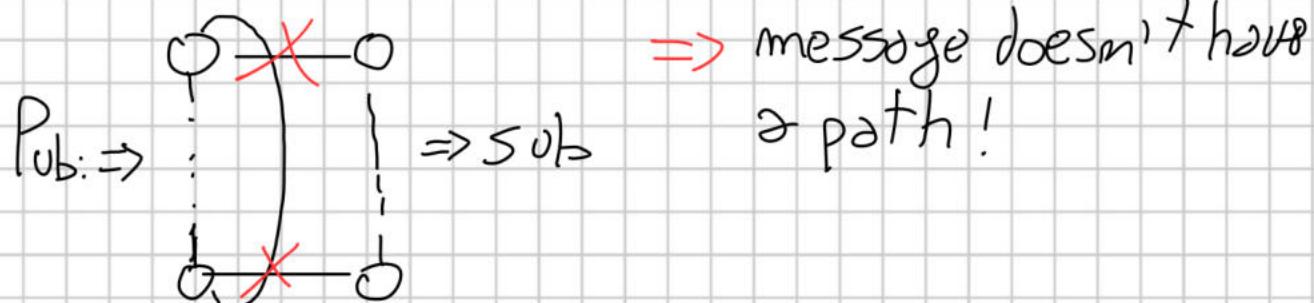
Upon event PP_{2P} Deliver (NOTIFY l, v) from P
if v \notin delivered do
 for each s \in subscription s.t. S.constraint(v) = True
 trigger PP_{2P} Send (PUBLISH l, v) to S.subscriber
 for each m \in neighbours s.t. m \neq p:
 trigger PP_{2P} Send (NOTIFY l, v) to m
 ↳ flooding of the event.
delivered = delivered \cup {v}

Upon EVENT PP_{2P} Deliver (PUBLISH l, v)
delivered = delivered \cup {v}
 ↳ subscriber save the value

Publisher with publish send his value to the k broker. The subscriber can subscribe/ unsubscribe sending the message to the k brokers. The brokers send the value received to a subscriber if exist a subscription saved and broadcast the value to the k broker near to him.

2. We can tolerate at most $k-1$ failures, because each time a failure, the neighbours of some brokers decrease, a partition is created, but until $k-1$ failures is impossible to have that a partition is isolated from an other, but with k failures can happen that a

partition is isolated and the NOTIFY events can't arrive to the subscribers:



At the beginning we have k partition of K brokers each, each partition have K connection to other partitions.

3. For tolerate the Byzantine processes we must use authenticated links, in each event we must check that the sender is equal to the sender in the event. Subscribers must checks the value received by brokers that is really a value on which is interested and that doesn't already received, also we can accept a value only when we have received $F+1$ events with the same value.

It is necessary that k is at least equal to $2F+1$ because we must guarantee a reliable communication.

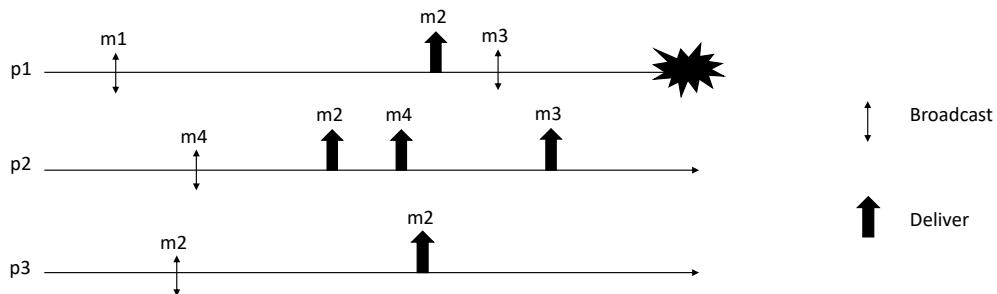
Distributed Systems (9 CFU)

05/07/2022

Family Name _____ Name _____ Student ID _____

Ex 1: Provide the specification of the (1, N) Regular Register and describe the majority voting algorithm discussed during the lectures.

Ex 2: Consider the message pattern shown in the Figure

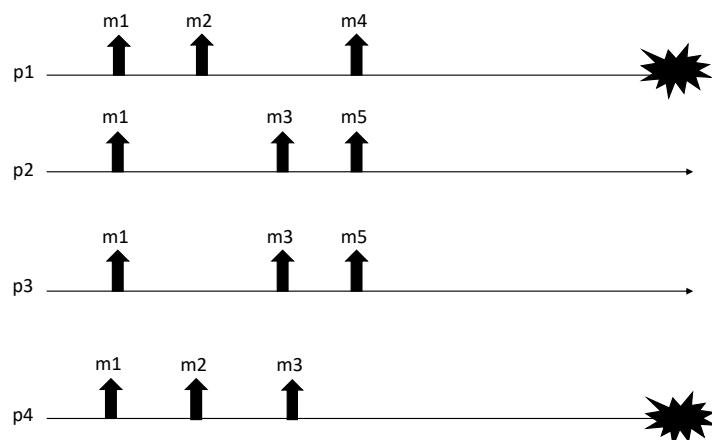


Answer to the following questions:

1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

NOTE: To solve the exercise you can just add deliveries of messages and new broadcast (if needed)

Ex 3: Consider the execution depicted in the Figure



Answer to the following questions:

1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO (UA, WUTO) but not TO (UA, SUTO)

3. Modify the run in order to obtain an execution satisfying TO (NUA, WNUTO) but not TO(NUA, WUTO).

NOTE: To solve the exercise you can just add deliveries of messages.

Ex 4: Let us consider a Regular Reliable Broadcast primitive satisfying the following properties:

- *Validity*: If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement*: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Let us consider a distributed system composed of N processes executing the Eager algorithm (reported in figure)

Algorithm 3.3: Eager Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb .

```

upon event (  $rb$ , Init ) do
   $delivered := \emptyset$ ;

upon event (  $rb$ , Broadcast |  $m$  ) do
  trigger (  $beb$ , Broadcast | [DATA, self,  $m$ ] );

upon event (  $beb$ , Deliver |  $p$ , [DATA,  $s$ ,  $m$ ] ) do
  if  $m \notin delivered$  then
     $delivered := delivered \cup \{m\}$ ;
    trigger (  $rb$ , Deliver |  $s$ ,  $m$  );
    trigger (  $beb$ , Broadcast | [DATA,  $s$ ,  $m$ ] );
  
```

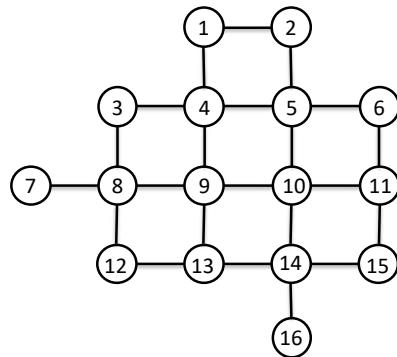
Answer to the following questions:

1. assuming that up to f processes may commit omission failures and no other failures may happen, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).
1. assuming that up to f processes may be Byzantine faulty but constrained to have a symmetric behaviour¹, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).

Ex 5: Consider a distributed system composed by n processes each one having a unique identifier. Processes communicate by exchanging messages through perfect point-to-point links and are connected through a grid (i.e., each process p_i can exchange messages only with processes located at *nord*, *sud*, *east* and *west* when they exist).

An example of such network is provided in the following figure:

¹ A Byzantine process has a symmetric behaviour if it can change the content of every message it is going to send, but it cannot send different values to different processes when invoking the $bebBroadcast$. Summarizing, it can cheat but it will do it in a consistent way.



Processes are not going to fail, and they initially know only the number of processes in the system N and the identifiers of their neighbors.

Processes in the system must agree on a color assignment satisfying the following specification:

Module

Name: k-Color assignment

Events:

Request: $\langle ca, \text{Propose} \mid c \rangle$: Proposes a color to be adopted.

Indication: $\langle ca, \text{Decide} \mid c \rangle$: Outputs a decided color to be adopted by the process

Properties:

Termination: Every process eventually decides a color.

Validity: If a process decides a color c , then c was proposed by some process or $c = \text{default}$.

Integrity: No process decides twice.

Weak Agreement: If two processes decide c_i and c_j then either $c_i = c_j$ or one of the two is default

Color Selection: Let C be the set of proposed colors, if $|C| > 0$ then there exists at least a color $c_i \in C$ that is decided by k processes.

Assuming that $1 < k < N$ and that k is known by every process, write the pseudo-code of an algorithm implementing the k-Color assignment primitive.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

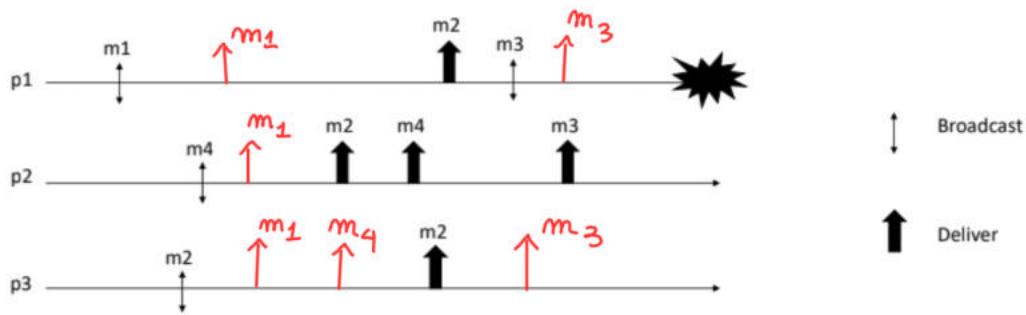
Signature: _____

Exercise 1.

(1, N) Regular Register specification. It has 4 events that are:
• $\langle \text{tt}, \text{Read} \rangle$ that invokes a read operation on the register.
• $\langle \text{tt}, \text{write } v \rangle$ that invokes a write operation with value v on the register.
• $\langle \text{tt}, \text{Read Return } v \rangle$ that completes a read operation on the register with return value v .
• $\langle \text{tt}, \text{Write Return} \rangle$ that completes a write operation on the register.
It must satisfy two properties: termination (if a correct process invokes an operation, then the operation eventually completes) and validity (A read that is not concurrent with a write returns the last value written; a read that is concurrent with a write returns the last value written or the value concurrently written). In the majority voting regular registers protocol we assumed a crash failure model with no perfect failure detector. We assumed N processes whose 1 writer and N readers, with a majority of correct processes. It uses P2P link and best-effort broadcast. The main idea of the algorithm is that each process locally stores a copy of the current value of the register. Each written value is univocally associated to a timestamp. The writer and the reader processes use a set of witness processes, to track the last value written. Use a quorum, the intersection of any 2 sets of witness processes is not empty for deliver a write or a read. The read return the highest value witness with the current timestamp when reach a quorum of reply and writer return is done when p reach a quorum of acks.

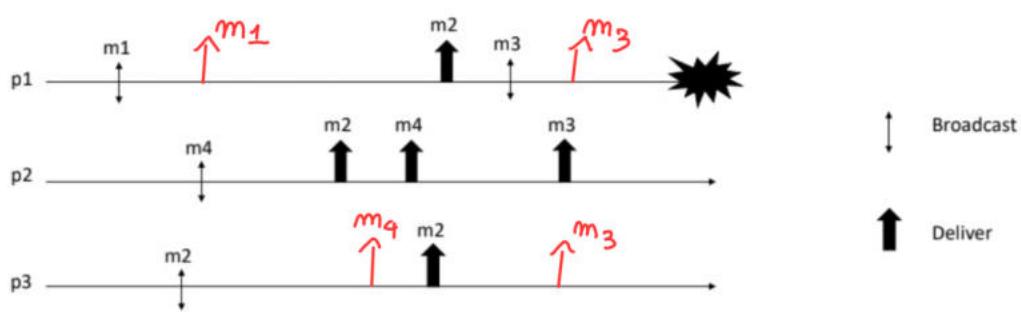
Exercise 2

1.



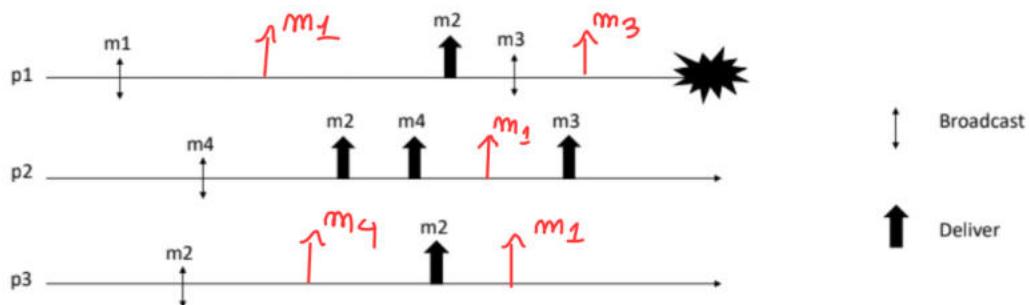
This turn satisfying Uniform Reliable Broadcast, because correct processes have same set of deliveries, the faulty one has a subset.

2.



This turn satisfying Regular Reliable Broadcast, because correct processes have same set of deliveries, but not uniform because faulty one have delivered m_1 that other don't have.

3.



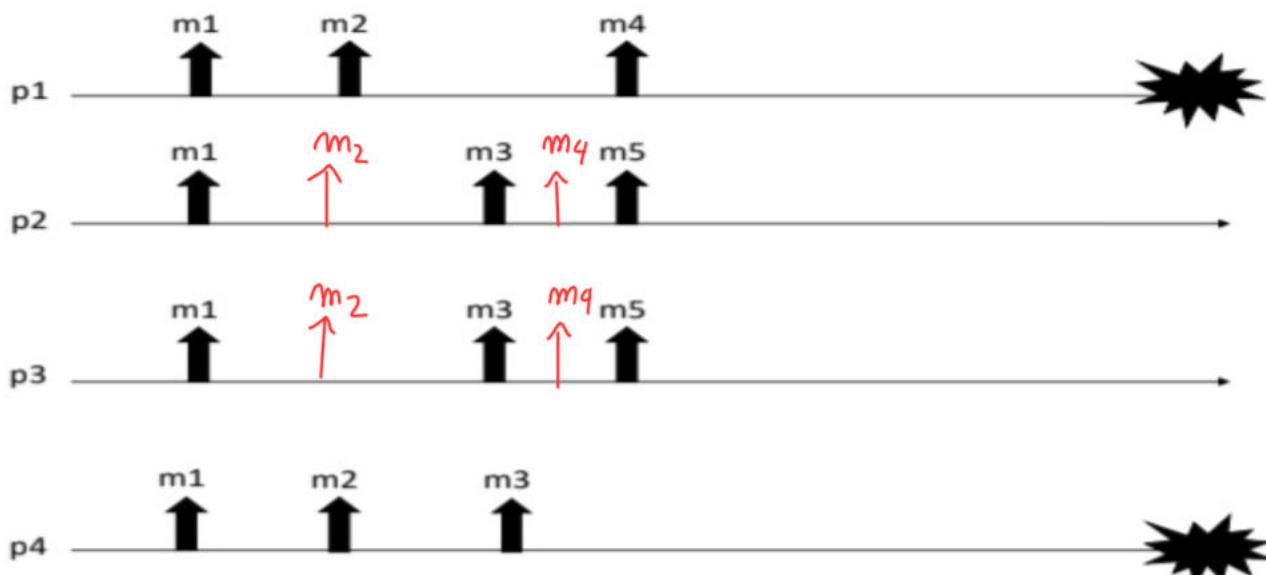
This turn satisfying best effort broadcast because correct processes don't have same set of deliveries.

4. For have causal order we must have $m_1 \rightarrow m_3$ and $m_2 \rightarrow m_3$, For total order because of P2 we have constraint $m_2 \rightarrow m_4 \rightarrow m_3$, the sequences that satisfying both are: (m_1, m_2, m_4, m_3), (m_2, m_1, m_4, m_3), (m_2, m_1, m_1, m_3), (m_2, m_1, m_1, m_3).

Exercise 3

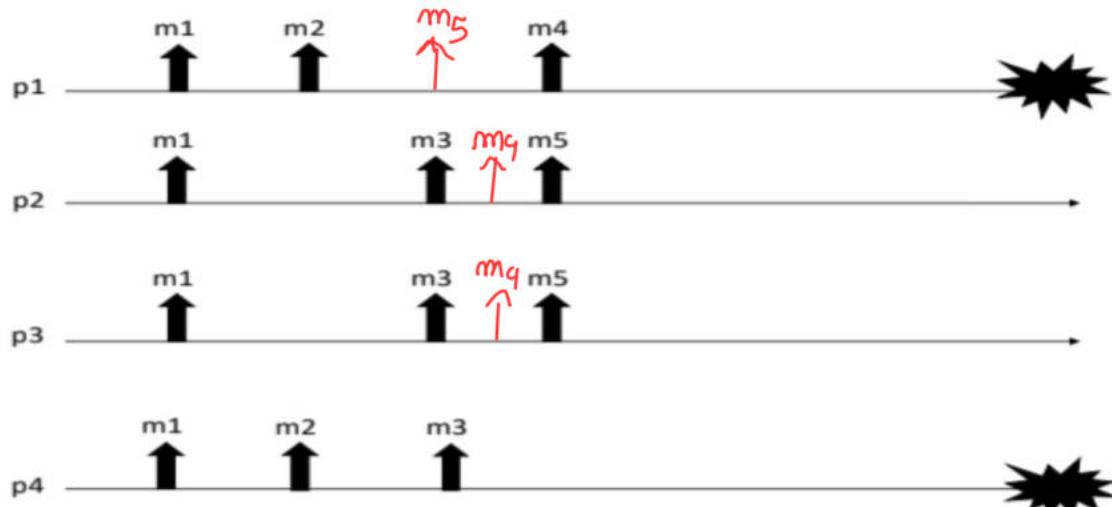
1. This turn doesn't satisfy UA because p_1 and p_4 , that are faulty, have not a subset of delivery of the correct processes. It satisfies SU TO because every process respect the constraint of mto, only pair are respected of other. TO (NUA, WUTO)

2.



it satisfies UA because correct processes have same set of messages, and faulty one a subset. It satisfies WUTO because the order of any pair is respected but not SU TO because p_1 deliver m_4 but not m_3 before

3.



it not satisfies UA because p_1 and p_4 deliver m_2 , it does not satisfy WUTO because p_2 violate order of pair $(m_1 \rightarrow m_3)$

Exercise 4

1. Validity: it is satisfied because if a correct process p broadcast a message m , eventually p receive $(p, l, DATA_S, m)$ and because $m \notin S$ delivered, he will deliver m . All work because p use P2P links.

No duplication: it is satisfied because when a process p receive a broadcast message m , check that $m \notin S$ delivered, it is impossible to have duplication.

No creation: it is satisfied because if m is delivered, then following the algorithm it must be that S generate a broadcast event with m , the F processes can only do an omission but not generate fake messages.

Agreement: it is satisfied because also if F processes faulty do an omission, when a correct process receive a message m do a broadcast, and other correct processes eventually deliver that message.

2. Validity: it is satisfied because if a correct process p broadcast a message m , eventually p receive $(p, l, DATA_S, m)$ and because $m \notin S$ delivered, he will deliver m . All work because p use P2P links.

No duplication: it is not guarantee because a byzantine process can not insert a message delivered in the set delivered, and to deliver the same message when receive again m .

No creation: it is not guarantee because a byzantine process can forge a message m with a fake sender, and because the algorithm don't use authenticate link, a process p will deliver this m .

Agreement: it is satisfied thanks to the fact we use P2P links and every correct process when receive a message m , he will generate a new broadcast of the same message, then at some point every correct process reach the same set of deliveries.

Exercise 5

We have m processes each with an unique identifier. We use P2P links, using a topology like in the figure, each process know and can communicate with north, south, east and west, if exist. No failures. Each process know only number of processes N . The algorithm that i propose do: is like a consensus protocol. Each process propose a color c ; and disseminate to his neighbours. When a process receive N different propose (each is a pair of (value, identifier)), it will decide for the darkest color (HP: colors are comparable). After he choose a value, given that we want that a color is decided by K processes, a process p check if its id is less or equal to K then generate a decide event with color c , otherwise it will decide default ($= \perp$).

next page pseudo-code:

two events and a passive function.

Init:

$mp = N$ (number of processes)

decision = (\perp , FALSE)

propose = \emptyset , where we collect proposes

neighbours = {nord, sud, east, west} (some not exist)

$K = *$ is given

Upon event $\langle c, \text{Propose} | c \rangle$ do

propose = propose $\cup \langle c, \text{self} \rangle$

for each $p \in \text{neighbours}$ do

trigger PP2PSend (IPROPOSE1, $\langle c, \text{self} \rangle$) to p

Upon event PP2PDeliver (IPROPOSE1, $\langle c, id \rangle$) from s

if $\langle c, id \rangle \notin \text{propose}$ do

propose = propose $\cup \langle c, id \rangle$

for each $p \in \text{neighbours}$ s.t. $p.i = s$ do

trigger PP2PSend (IPROPOSE1, $\langle c, \text{self} \rangle$) to p

else

nothing i have already received

Upon exist $|\text{Propose}| == N$ and $\text{decision}_2 == \text{FALSE}$ do

choose = darkest_color (Propose) \downarrow

If self $\leq K$ do

decision = (choose, TRUE)

return darkest color proposed

trigger $\langle c, \text{decide} | \text{choose} \rangle$

else

decision = (\perp , TRUE)

trigger $\langle c, \text{decide} | \perp \rangle$

it satisfies termination, because every process disseminate his propose and the received one to neighbours, given that there are no failures and P2P links are used, all processes will receive N different proposal and decide a value. It satisfies validity because when a process decide, it decide the darkest color proposed by a process or the default. it satisfies integrity because when a process decide, set decision = 2 equal to TRUE and the passive check is no longer satisfied. It satisfies weak agreement because no process decide a value different from darkest color of 1, darkest color is the same for all because all processes have same set propose. It satisfies color selection because we construct algorithm in a way that the darkest color is chosen by the first K processes, exist always a color chosen k times.